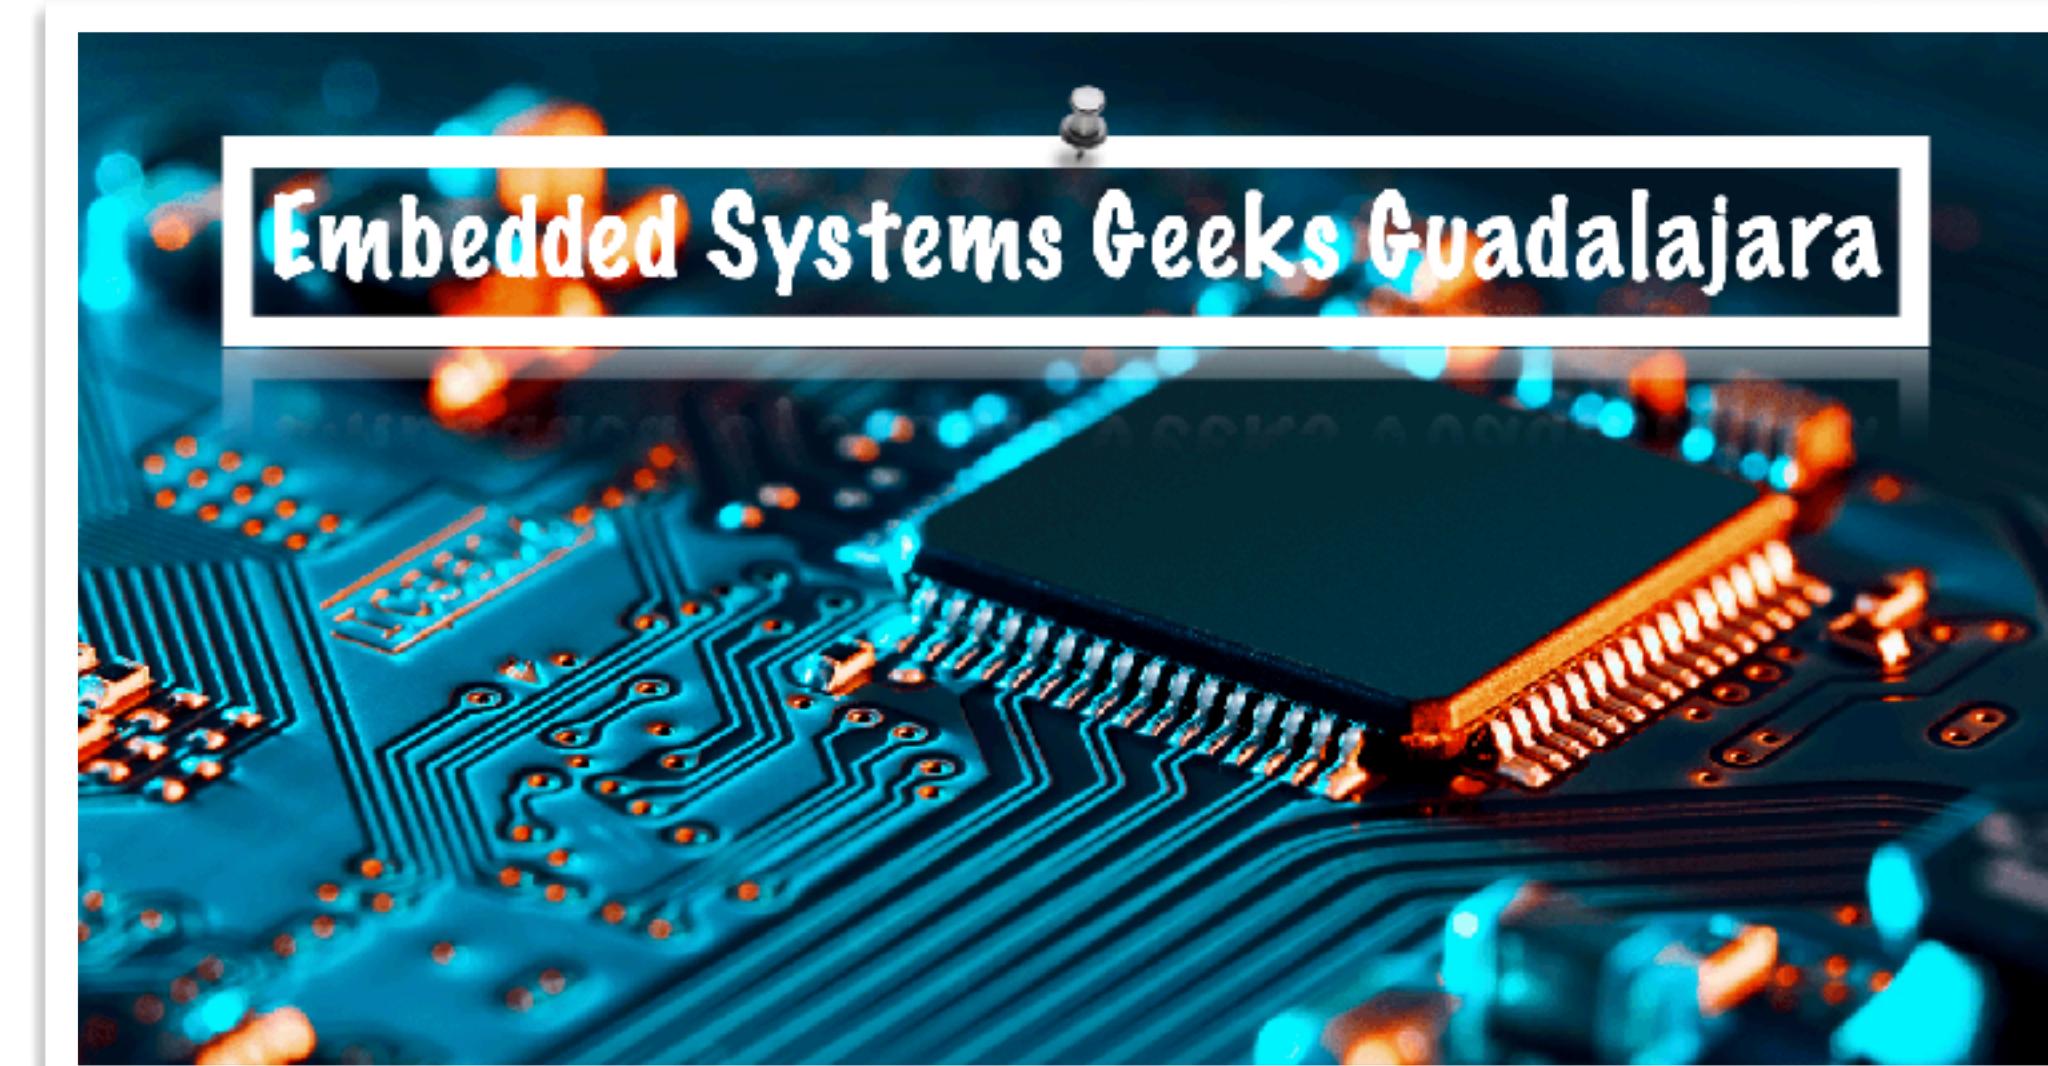


Embedded Systems Geeks GDL

Arduino & AVR: a deep(er) dive

Dr Frank Zeyda - 19 September 2023 @ HackerGarage GDL



Frank Zeyda

Dipl.-Inform. BSc PhD PgCLTHE

Independent Consultant
Safety-Critical Systems

<https://www.linkedin.com/in/frank-zeyda/>



Research Interests:

- Semantic Foundations / Unifying Theories (UTP)
- Formal Methods and Verification
- Rigorous Digital Engineering
- Automated Theorem Proving (Isabelle/HOL)



frank.zeyda@gmail.com

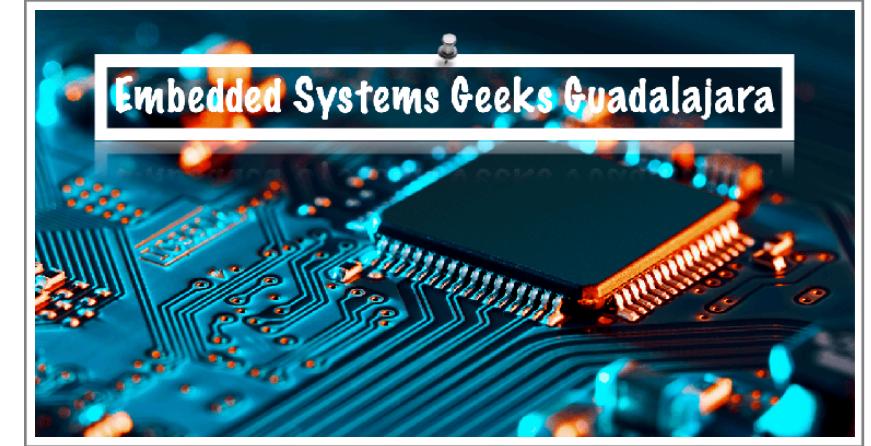
Academic Involvement (2001 to 2018):

- PhD from Teesside University (UK) in 2007.
- Research Associate/Fellow at the University of York (UK).
- Senior Lecturer at Teesside University and the University of York (UK).
- Lead RA on several EPSRC-funded and EC-funded research projects.
- Most publications are listed on: <https://dblp.org/pid/23/6727.html>

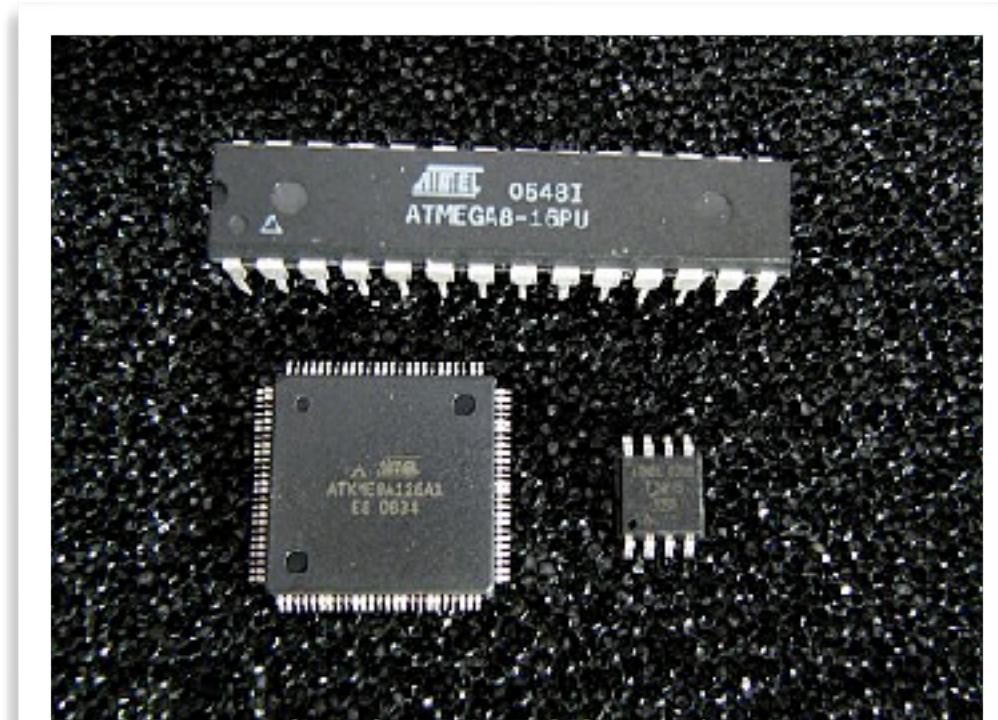
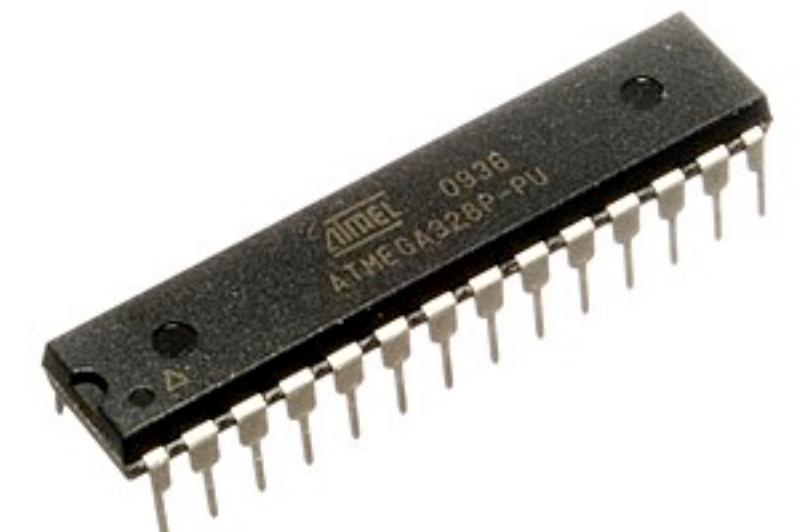
Industrial Involvement (1999 to 2000 & 2018 to 2022):

- Siemens Mobility / Transportation Systems (Germany)
- Verified Systems International GmbH (Germany)
- Galois, Inc. (U.S.) (consultant on industrial R&D projects)

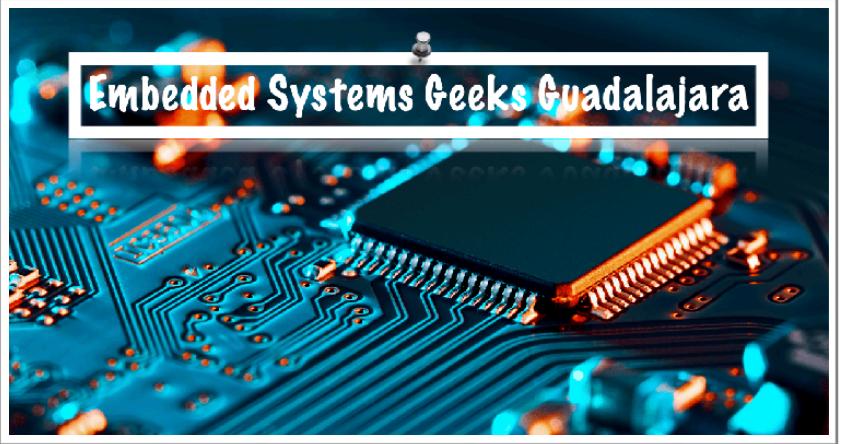
Introduction



- Topic of this talk: learn more about the **AVR architecture**
- **Arduino**: name of the “**brand**” and prototyping boards
- **AVR**: architecture of the underlying microcontroller family
- Commonly accepted AVR stands for **A**lf and **V**egard's **R**ISC processor, although Atmel denies it is an acronym (conceived by **two students**)
- Original AVR MCU developed by Nordic VLSI / Semiconductors (Norway) and later sold to Atmel Norway (subsidiary of Atmel).
- First AVR microcontrollers (8 bit) emerged in 1997.
- Atmel had shipped **500 million** AVR flash microcontrollers.

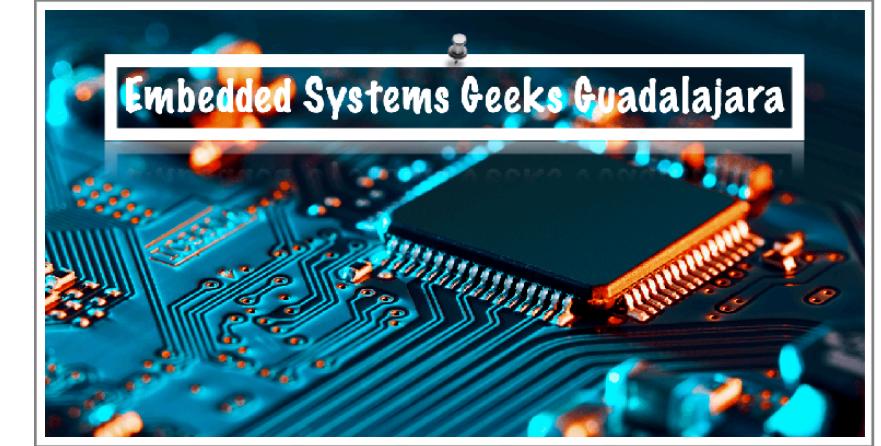


Overview of the talk



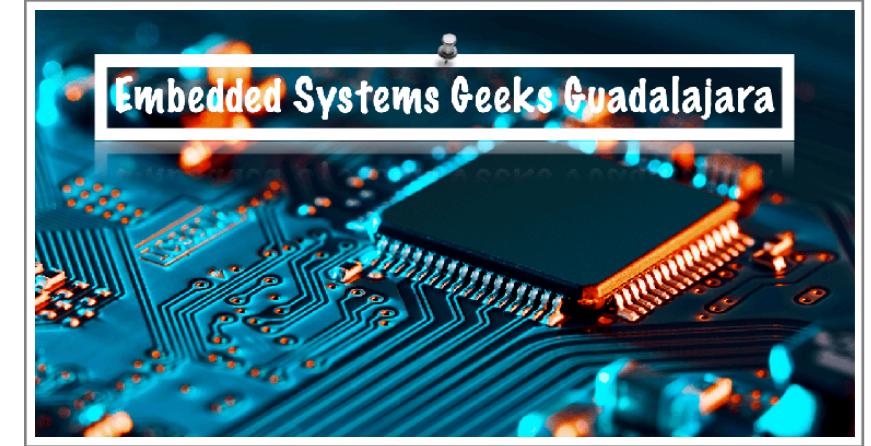
- Arduino & AVR Architecture 101
- Knowing your Chip
- **Example:** OLED Device Driver
- Communication Protocols
- Miscellaneous Topics
- Conclusion & Further Reading

AVR Architecture 101



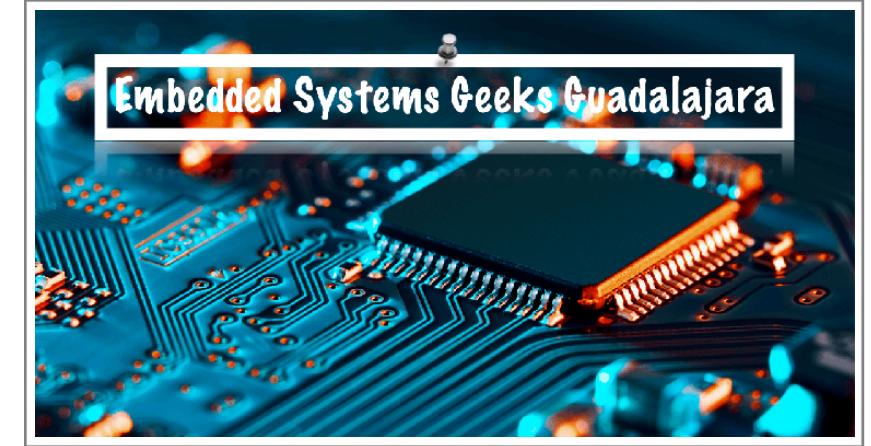
- Modified **Harvard** architecture: programs and data are stored in separate memories.
- Compare this to **von Neumann** architecture, where **all memory** is capable of storing program instructions
- Example: ATmega328P (Arduino Uno) has ...
 - ➔ 32 KB of ISP flash memory for programs, but only
 - ➔ 2 KB of static RAM (SRAM) for dynamic data
- Flash memory is usually **protected** during program executions (**why?**) and not meant to be written.
 - ◆ But, apparently, this may *still* be possible with a custom boot loader ...

Let's talk more about memory



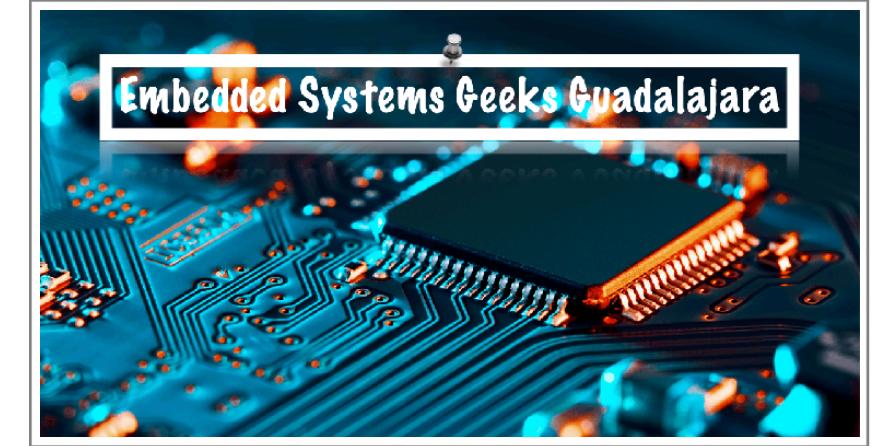
- When doing **Desktop / Cloud** application development, we usually do not think (very much) about memory (my MacBook has **64 GB**).
- Perhaps those developing **web** or **mobile** apps are more tuning into this.
- But this is still nothing compared to the restrictions on some MCUs.
- **Problem:** How to work with:
 - ▶ **32 KB** of program space (**flash, read-only**)
 - ▶ **2 KB** of (S)RAM for data (**volatile, read/write**)
 - ▶ **1 KB** of EEPROM (**non-volatile, read/write**)
- **Solution:** We have to be creative!

Let's talk more about memory



- Memory utilisation economy might imply ...
 - ✓ Strip dead **unused functions / code** from executable
 - ✓ Only put into **SRAM** what needs to be there ... (e.g., constant data can be in program memory)
 - ✓ Beware of using **large data structures** (arrays, look-up tables, etc.)
 - ➡ E.g., ask yourself: do we **really need** a frame buffer?
 - ✓ Strip all debug symbols and debug information (no gcc -g ...)
 - ✓ Code thoughtful and efficiently (modularisation, optimization, reuse)
 - ➡ What matters is the size of the binary, not the source code.

Let's talk more about memory

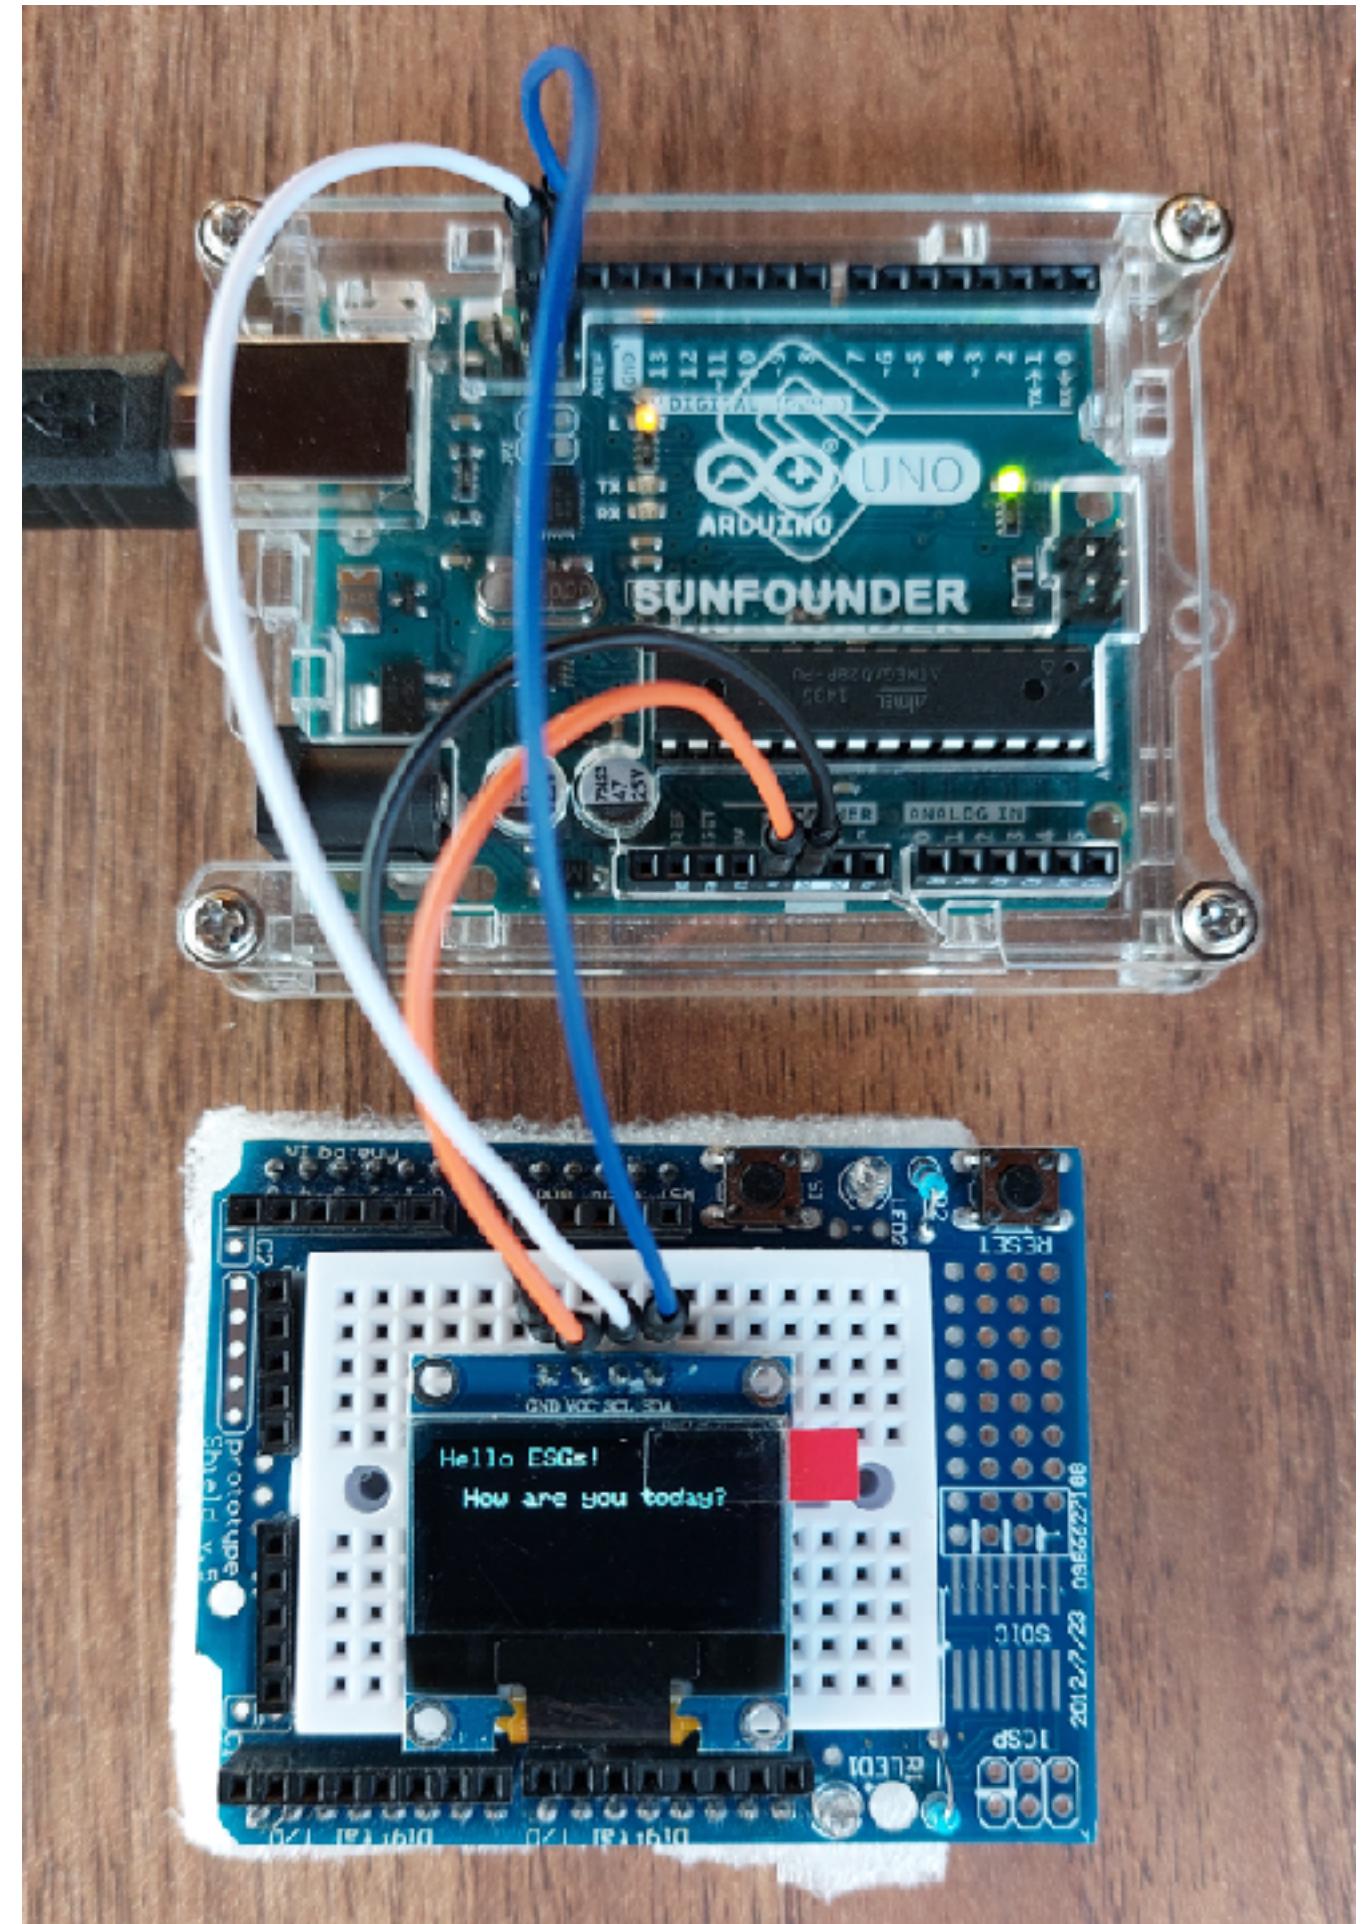


- Some personal experiences:
 - Third-party and open-source code is often needlessly bloated.
 - **Example:** Adafruit GFX library takes away 1 KB of SRAM for a frame buffer, i.e., when using a 128x64 pixel OLED display (display needs **no refresh**).
- Sometimes, ready-made components do **more than needed**
 - **Example:** full **I²C compliance** must support **multiple masters & slaves**, **arbitration**, **clock stretching**, **read/write**, etc.
 - But if there exists only **one** master & slave with a simple one-directional communication pattern, we do not need all of that ...
- Recursion can be **dangerous** due to the limited stack size.
- We often have a **trade-off** between **speed** and (memory) **resource efficiency**.

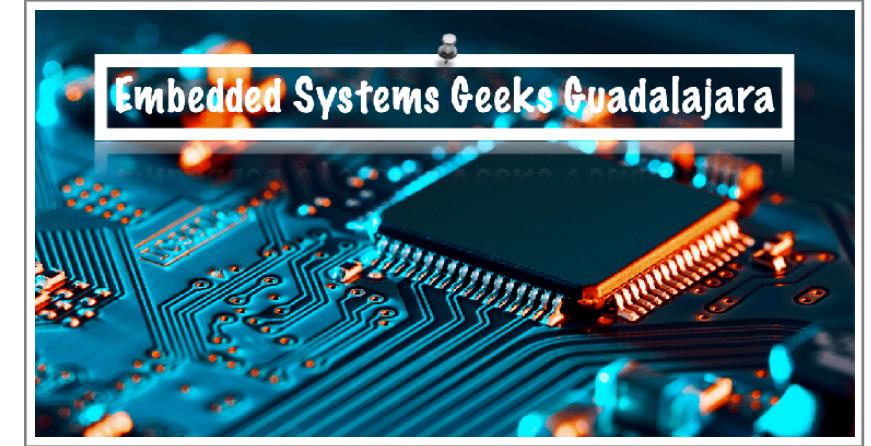
Example: Not using flash memory



```
1 #include "oled.hpp"
2
3 /* Implicitly, all constants are put into SRAM. */
4 static const char s1[] = "Hello ESGs!";
5 static const char s2[] = "How are you today?";
6
7 /* The setup() function is called only once by the AVR run-time. */
8 void setup() {
9     oled::init();
10    oled::clear();
11 }
12
13 /* The loop() function is called repetitively by the AVR run-time. */
14 void loop() {
15     oled::move(0, 0);
16     oled::print(s1);
17     oled::move(2, 10);
18     oled::print(s2);
19     while (true) { /* wait forever */ }
20 }
```



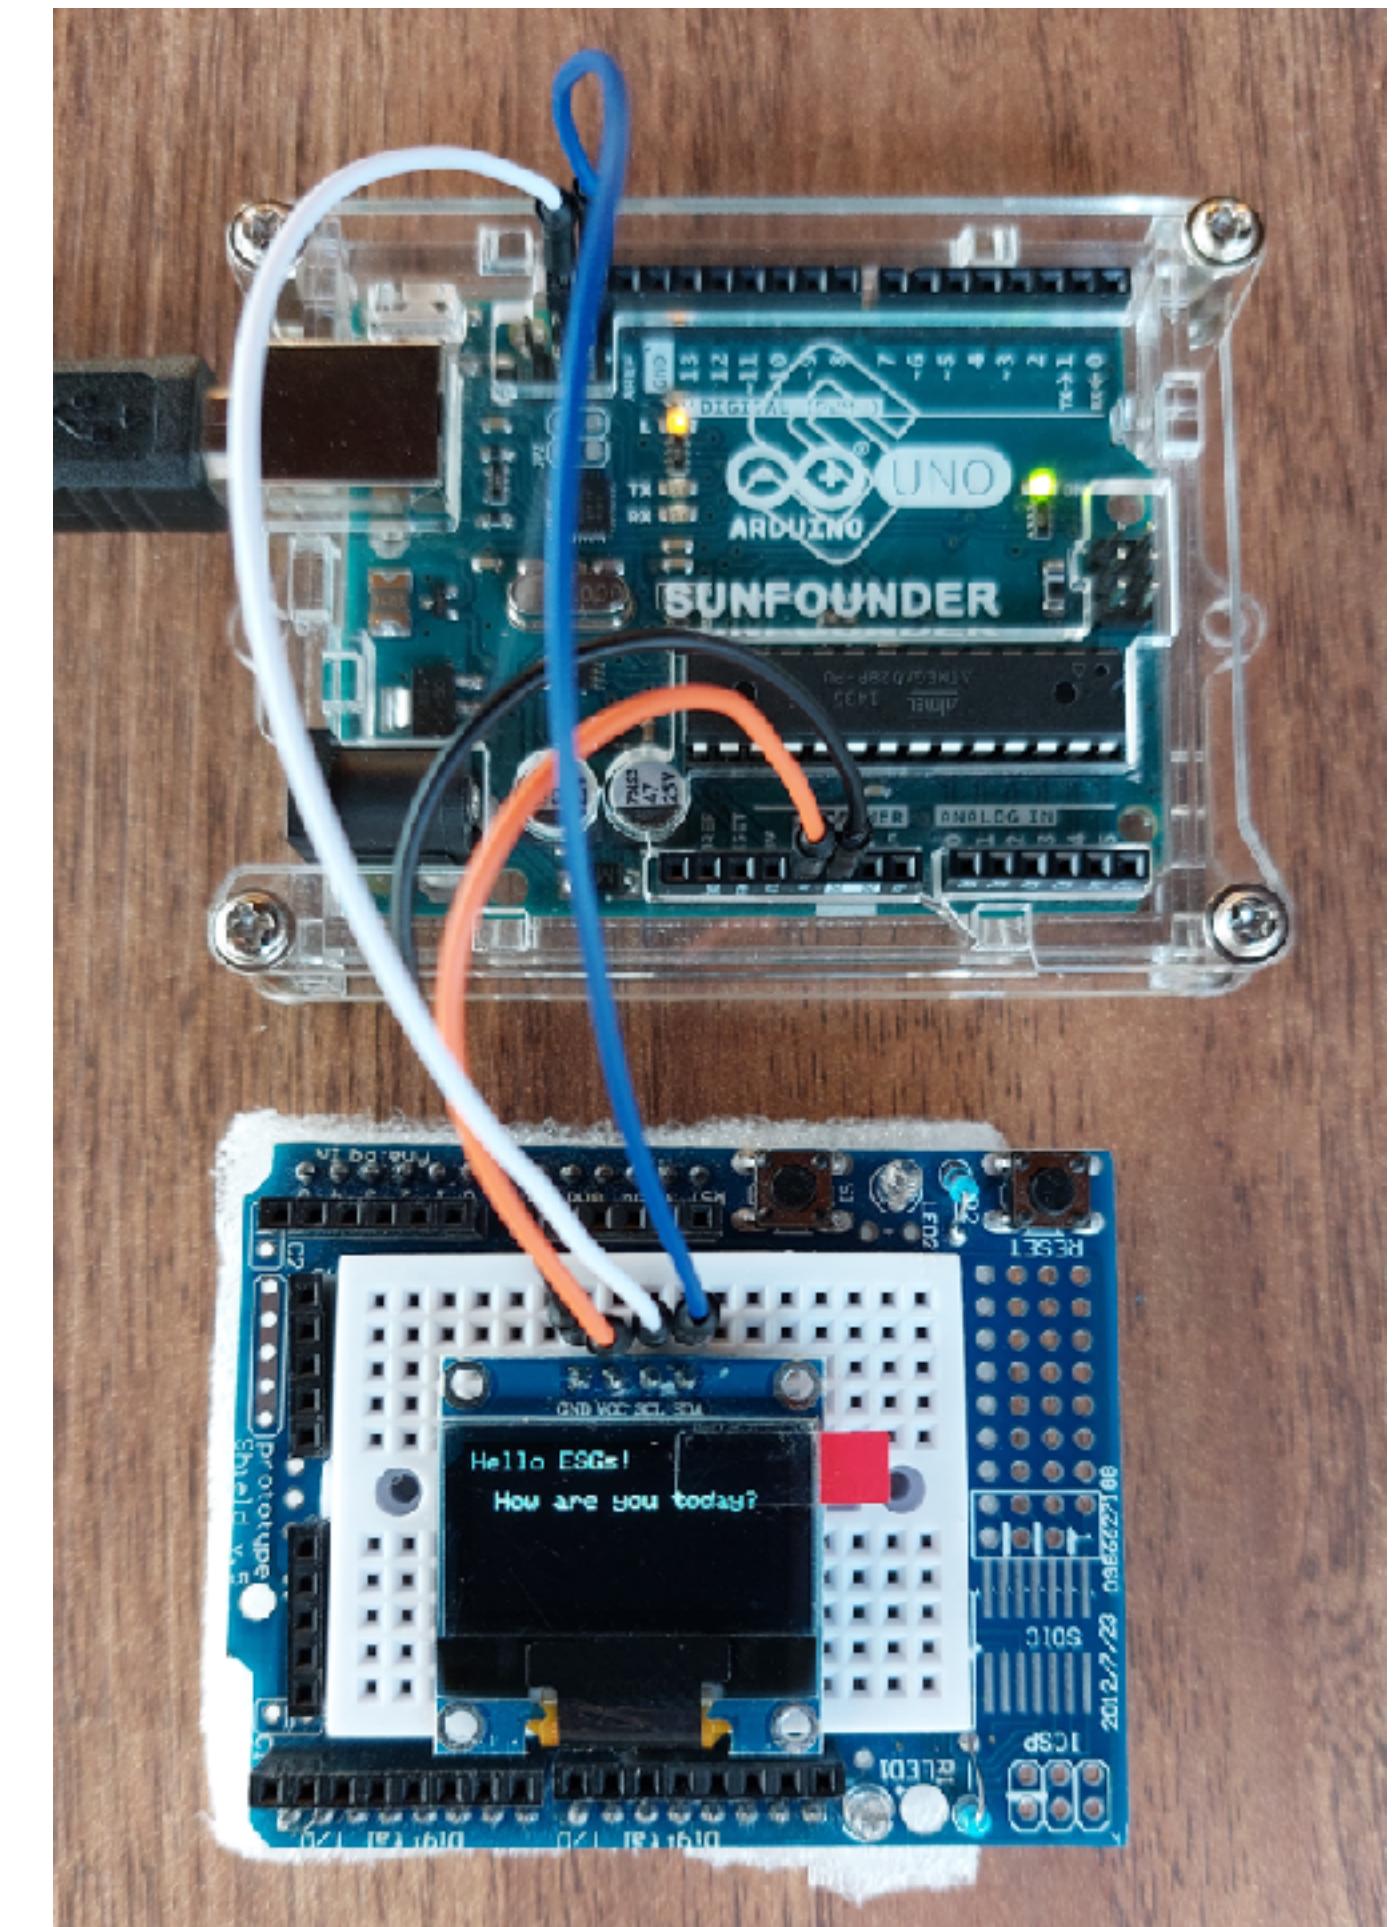
Example: Not using flash memory



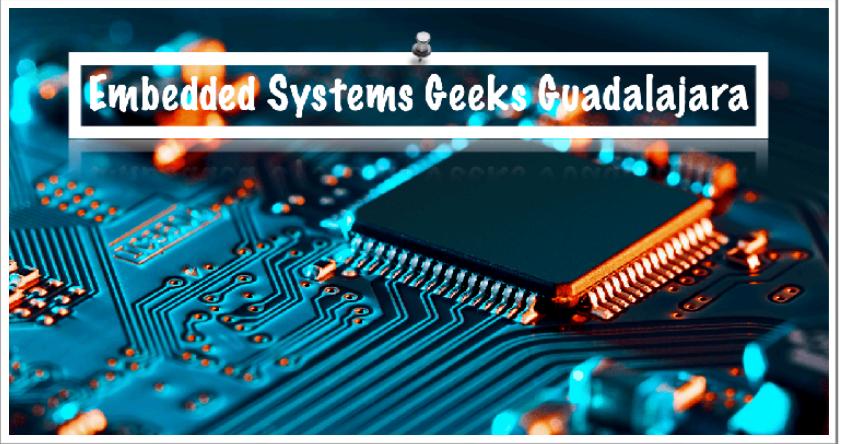
```
1 #include "oled.hpp"
2
3 /* Using the below with oled::print_p(...) saves precious SRAM. */
4 static const char PROGMEM s1[] = "Hello ESGs!";
5 static const char PROGMEM s2[] = "How are you today?";
6
7 /* The setup() function is called only once by the AVR run-time. */
8 void setup() {
9     oled::init();
10    oled::clear();
11 }
12
13 /* The loop() function is called repetitively by the AVR run-time. */
14 void loop() {
15     oled::move(0, 0);
16     oled::print_p(s1);
17     oled::move(2, 10);
18     oled::print_p(s2);
19     while (true) { /* wait forever */ }
20 }
```

Add attribute for linker to place data into flash

Custom print function that read strings from program/flash memory

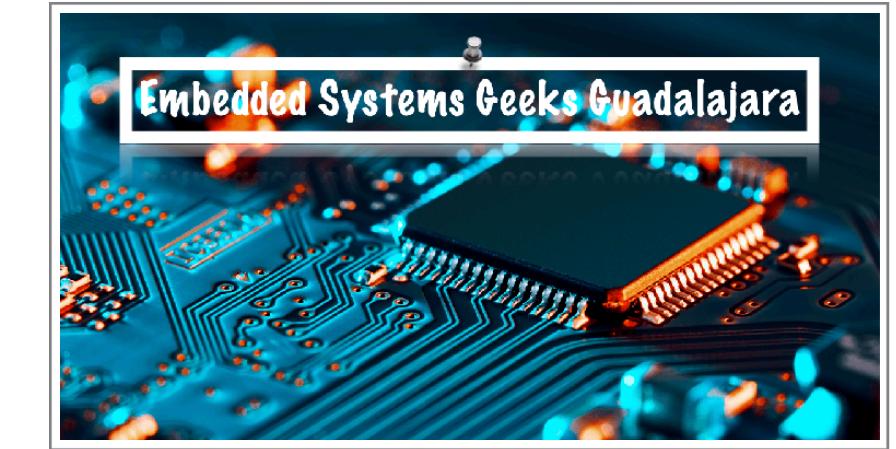


Difference in Compilation Output



- **TODO:** Show difference in compilation output here!

Accessing Program Memory (API)

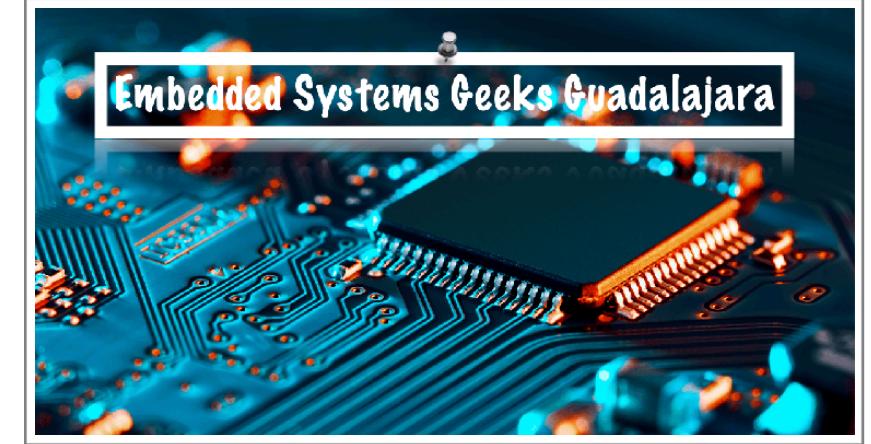


- Side-by-side comparison of `print(...)` and `print_p(...)`:

```
uint8_t print(const char* s) {
    if (s != nullptr) {
        while (*s != '\0') {
            const uint8_t error = emit(*(s++));
            if (error != 0) { return error; }
        }
    }
    return 0; // no error if reaching here
}
```

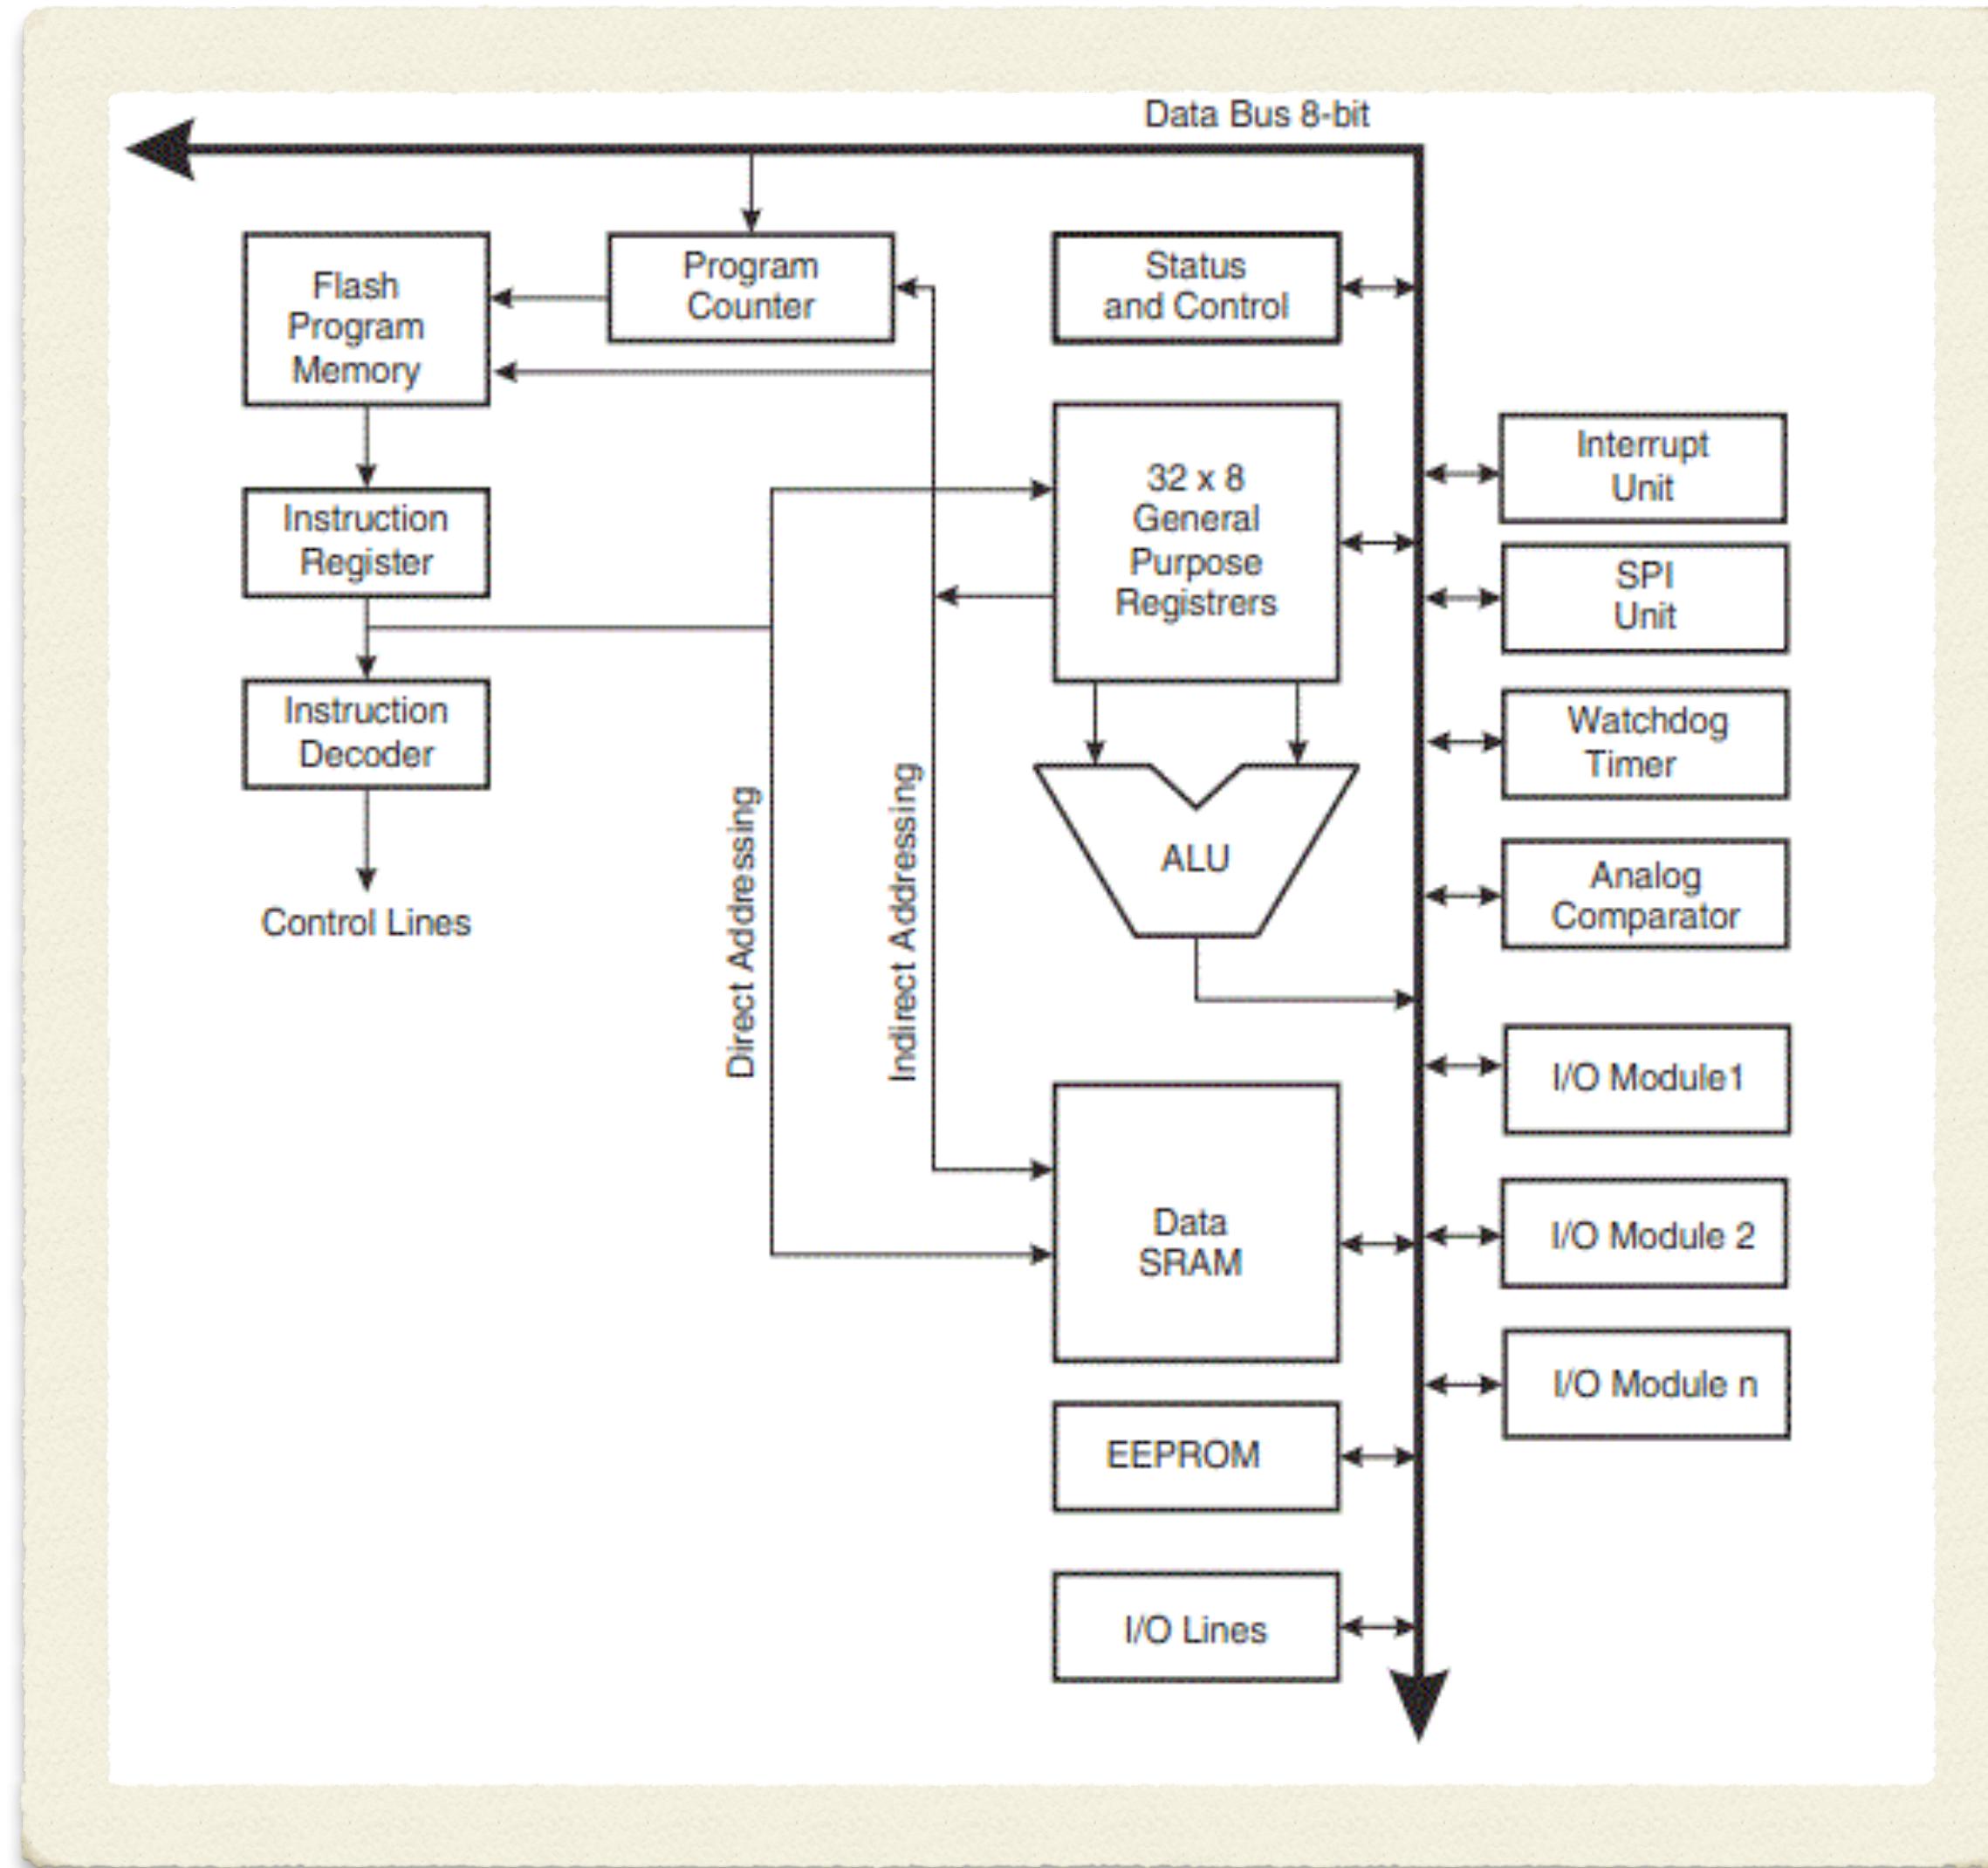
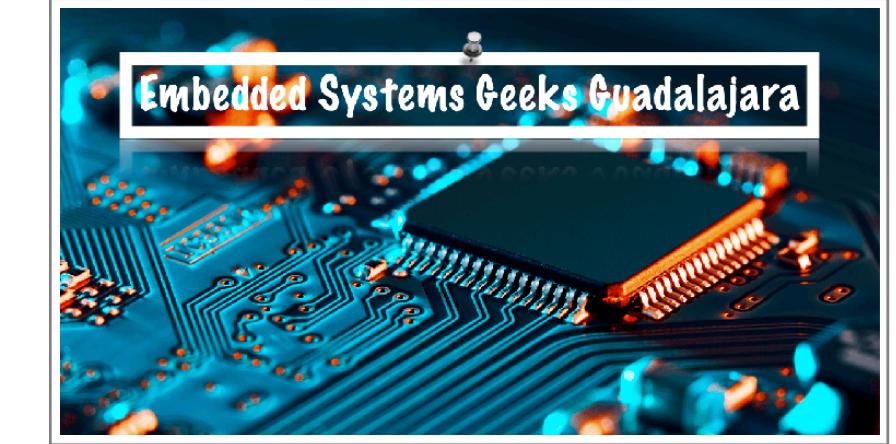
```
uint8_t print_p(const char* s) {
    if (s != nullptr) {
        char c = pgm_read_byte(s++);
        while (c != '\0') {
            const uint8_t error = emit(c);
            if (error != 0) { return error; }
            c = pgm_read_byte(s++);
        }
    }
    return 0; // no error if reaching here
}
```

Memory on the Atari 2600



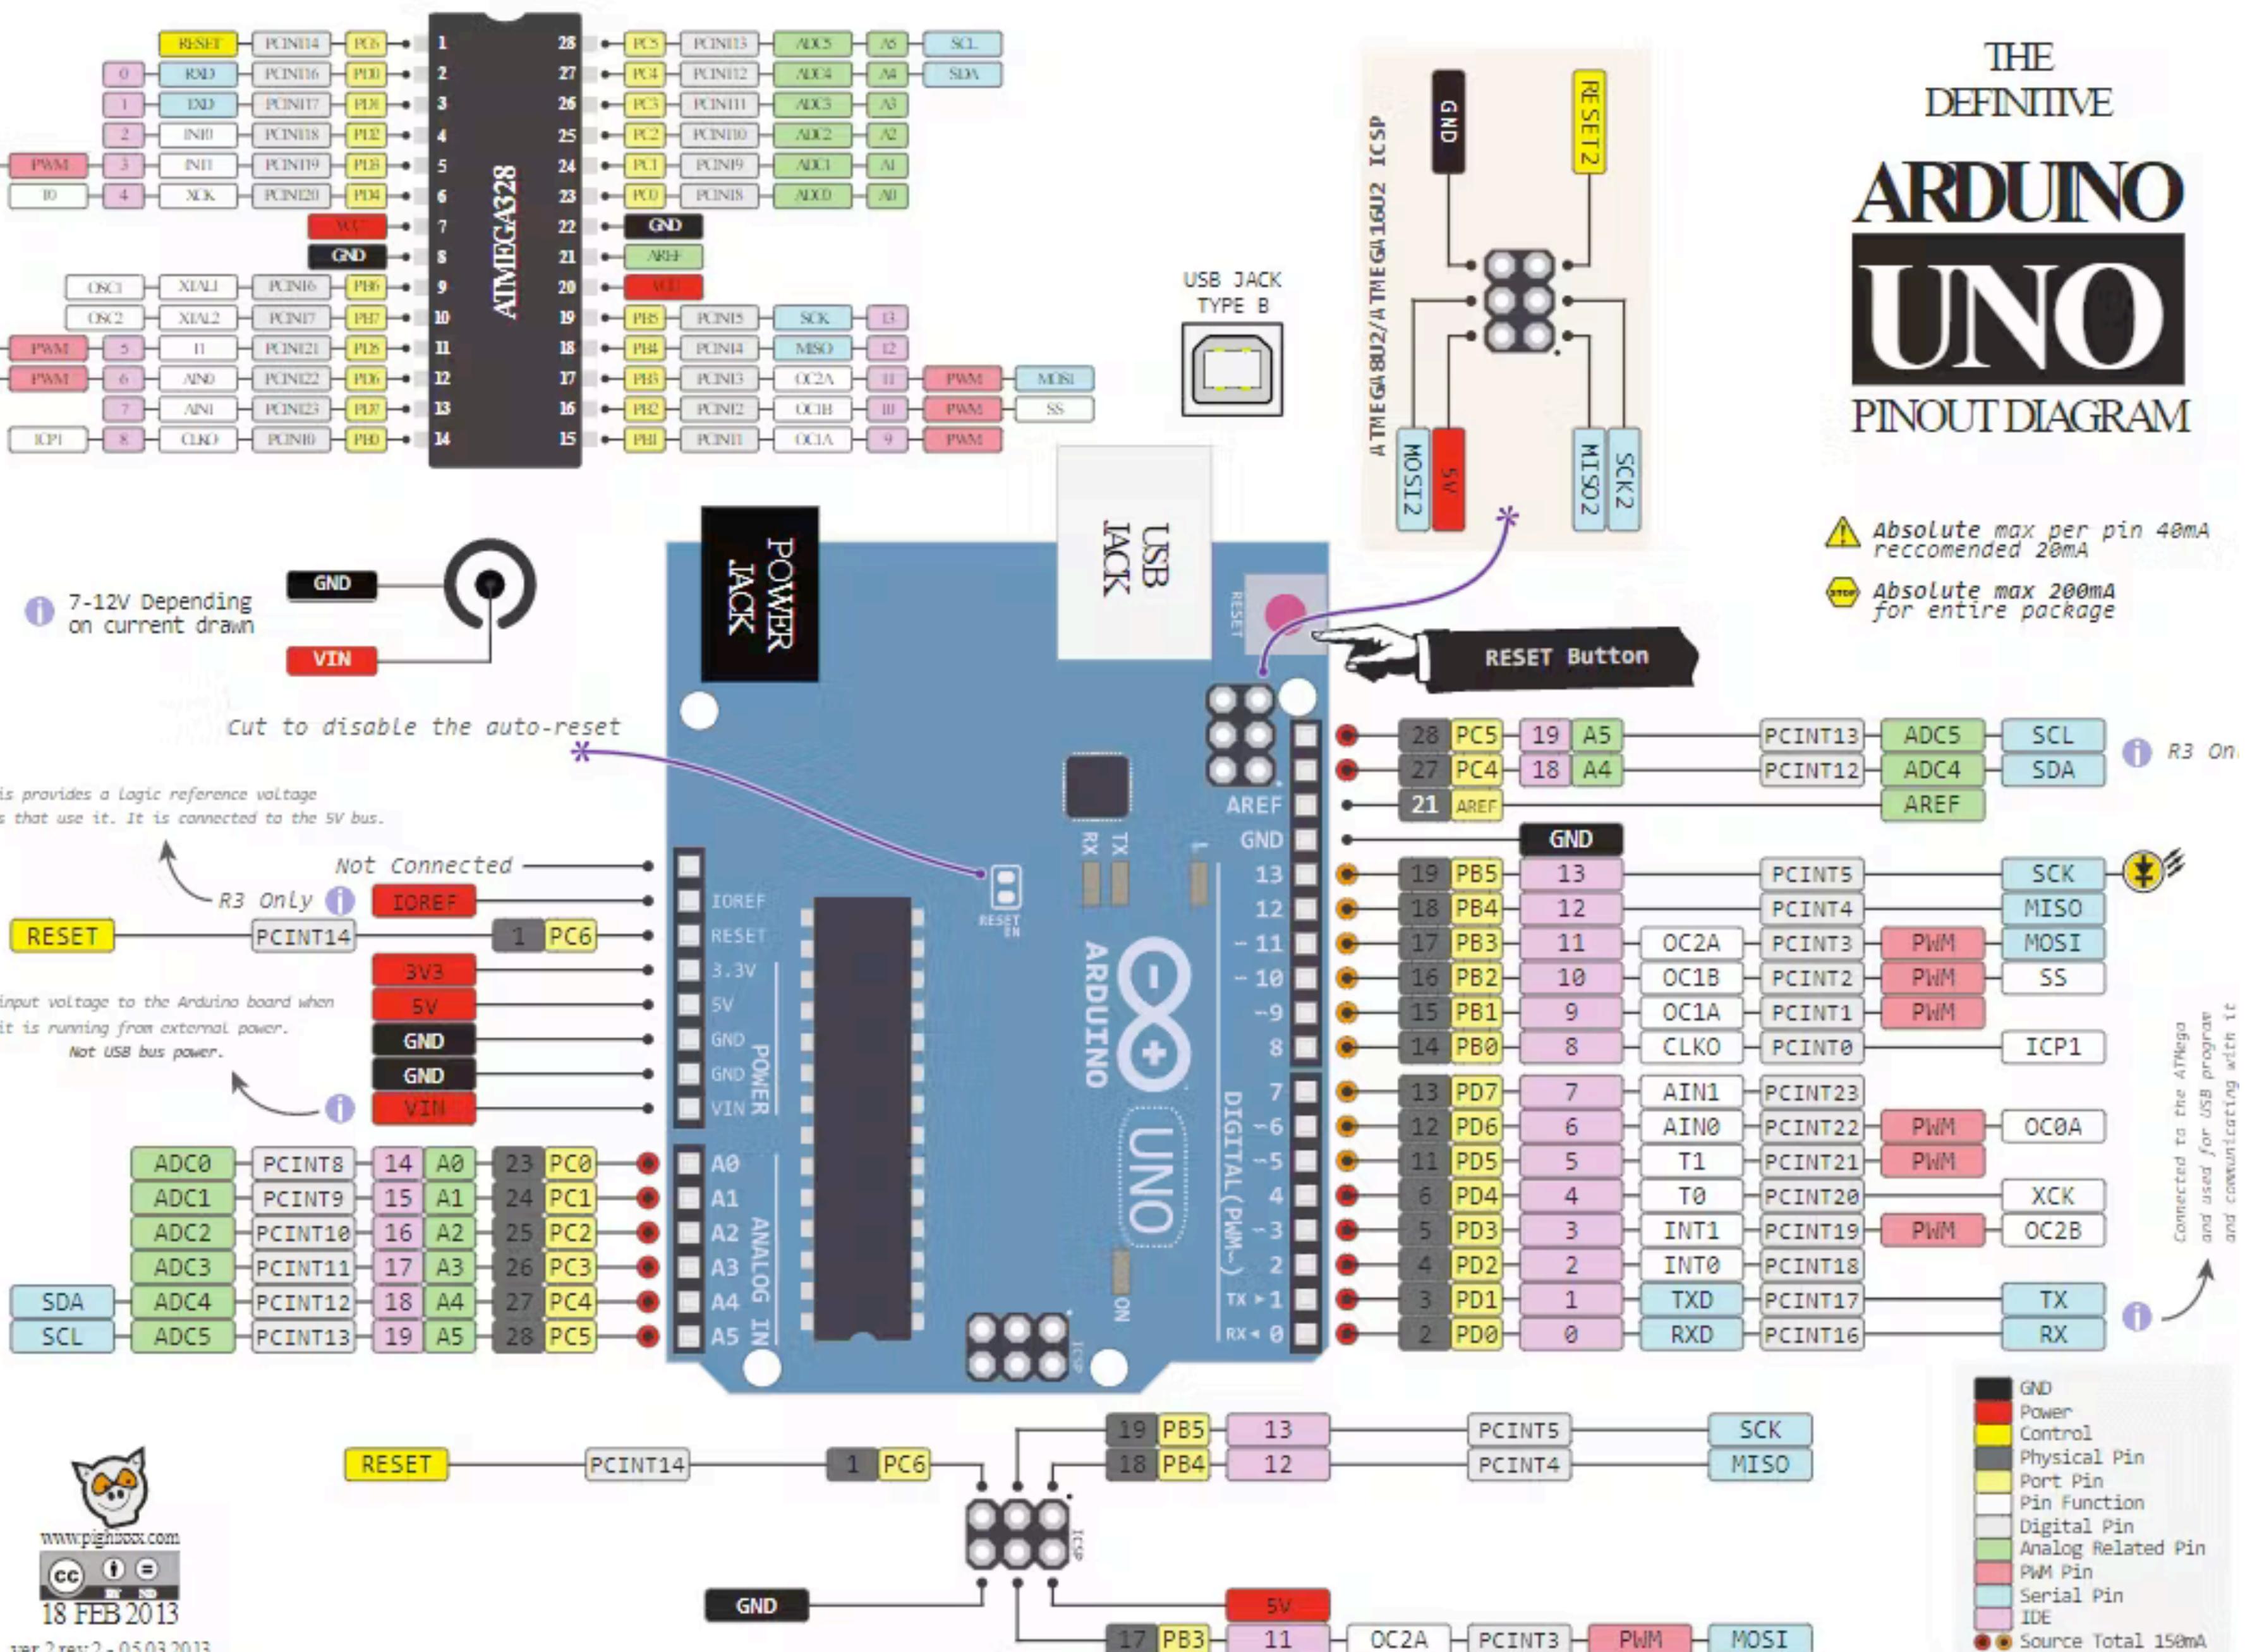
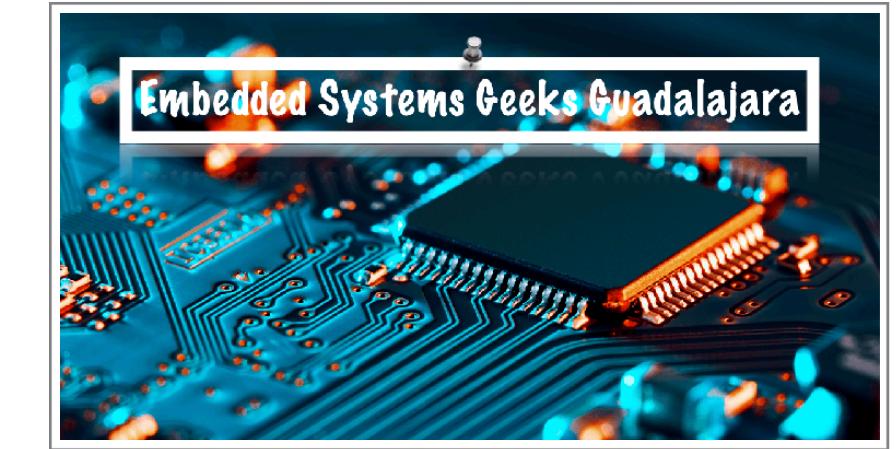
- But consider yourself lucky. On the Atari 2600 ...
 - ... we only had **128 bytes** of writable memory
 - ... most early games had to fit into 4 KB of program memory
 - ... the console had **no frame buffer** — the CPU had to write data into the **Audio & Video Processor (TIA)** during HW scan-line generation.
 - ... this leaves limited time to run **game logic (mostly V/H blank)**
 - ... and there was no **DMA** and surely no **GPU** or **coprocessors** ...
- Still developers managed to create some cool games for the system.

ATmega328 Internals

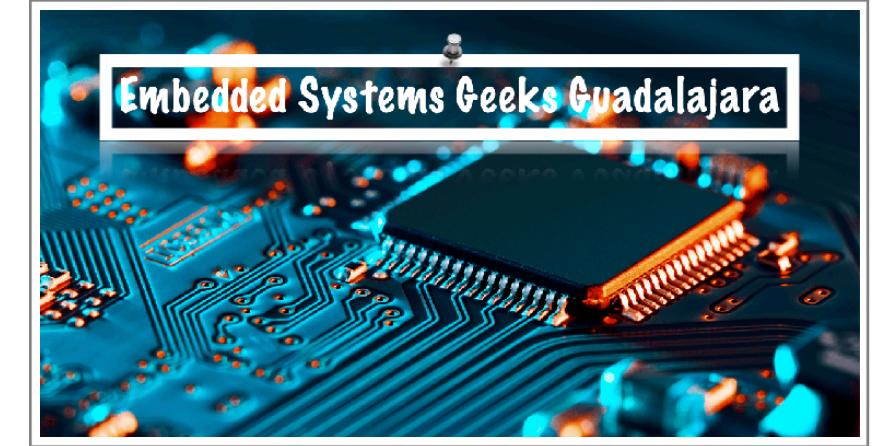


- 32 General Purpose Registers
 - 8-bit data width
- 20 MHz clock frequency
- Interrupt Control Unit (**ICU**)
- 8-bit / 16-bit Timers
- Programmable serial **U(S)ART**
- Master/slave **SPI** serial interface
- Byte-oriented 2-wire serial interface
 - Phillips **I²C** compatible
- **PWM** & 10-bit **ADC** channels (six)
- 23 programmable I/O lines
- Watchdog Timer

THE
DEFINITIVE
ARDUINO
UNO
PINOUT DIAGRAM

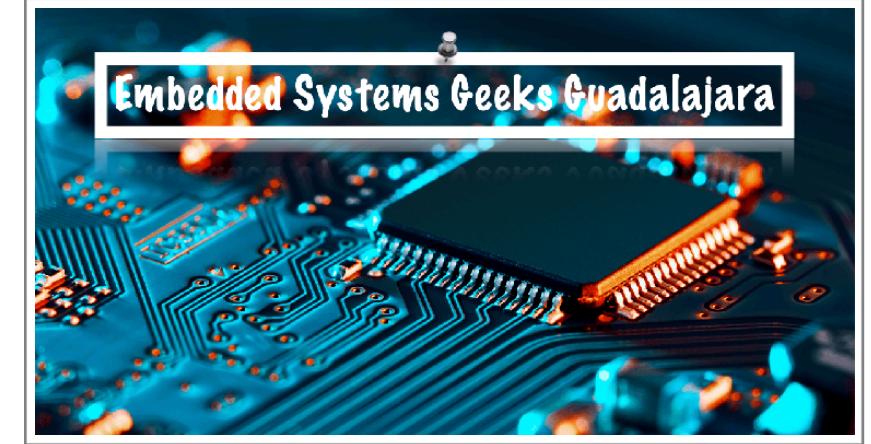


The Anatomy of an MCU



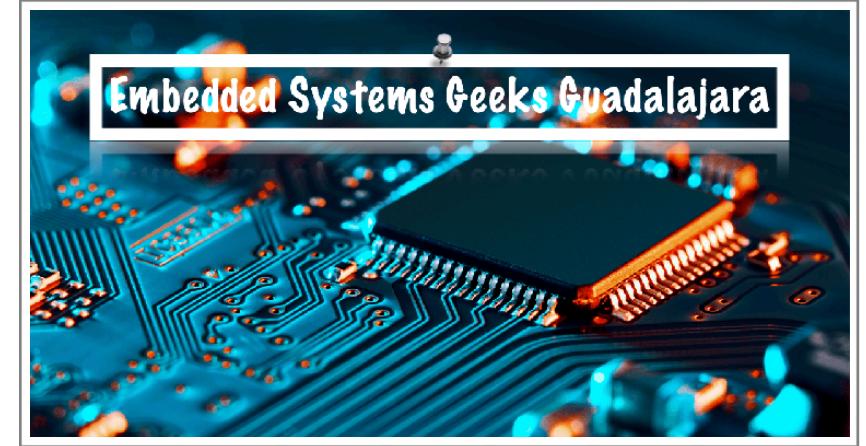
- We already learned that MCUs can be resource-constrained:
 - The AVR family is also includes larger memory MCUs (**megaAVR**, etc.)
- On the other hand, many MCUs come with a lot of useful embedded components to communicate with **other** hardware devices:
 - Natively speaking UART, SPI, I²C, even [IEEE 802.11 WiFi](#)
 - See this [LINK](#) for an useful set of slides on this topic.
- The **RP2040** MCU (Raspberry Pico) even provides state-machine coprocessors to implement your own low-latency protocols (e.g. **VGA**) ...
- This enables us to trade **bit-banging** for **byte-banging** ...
- Together with DMA, this causes next-to-no CPU load (**no busy waiting**)

Knowing your Chip

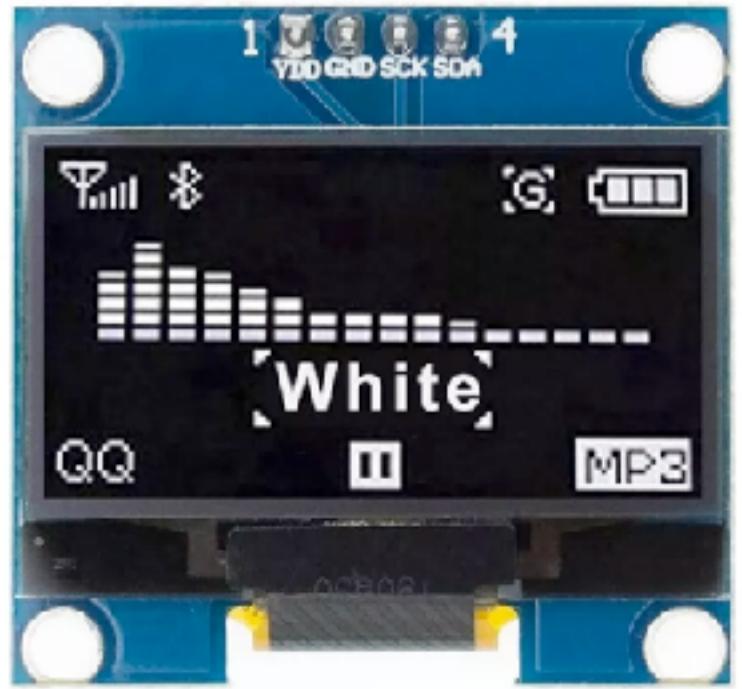


- Application developers commonly have to study **API** and **Framework** documents
- Hardware engineers are used to spend a lot of time reading **data sheets**.
- These can be intimidating at first sight ...
 - Often go into **thousands** of pages (**Pepe?**).
- Data sheets are almost **never read** from start to finish ...
 - Typically, you just pick the section that is relevant to your problem!
- Data sheets vary in terms of **quality**, **completeness**, **usability**, etc.
 - Do not despair. Sometimes, **wikis**, **blogs**, and **forums** offer more help.
- Below a link to the data sheet for the **ATmega328** (**294 pages**):
[https://ww1.microchip.com/downloads/en/DeviceDoc/Atmel-7810-Automotive-Microcontrollers-ATmega328P Datasheet.pdf](https://ww1.microchip.com/downloads/en/DeviceDoc/Atmel-7810-Automotive-Microcontrollers-ATmega328P_Datasheet.pdf)

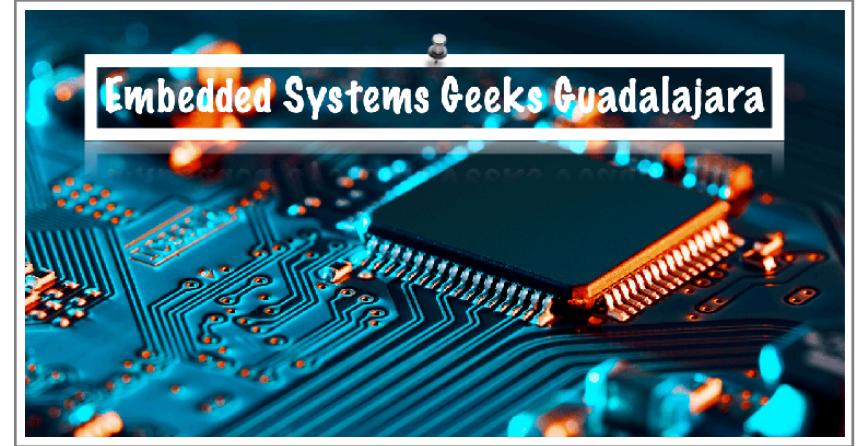
Knowing your Chip (cont'd)



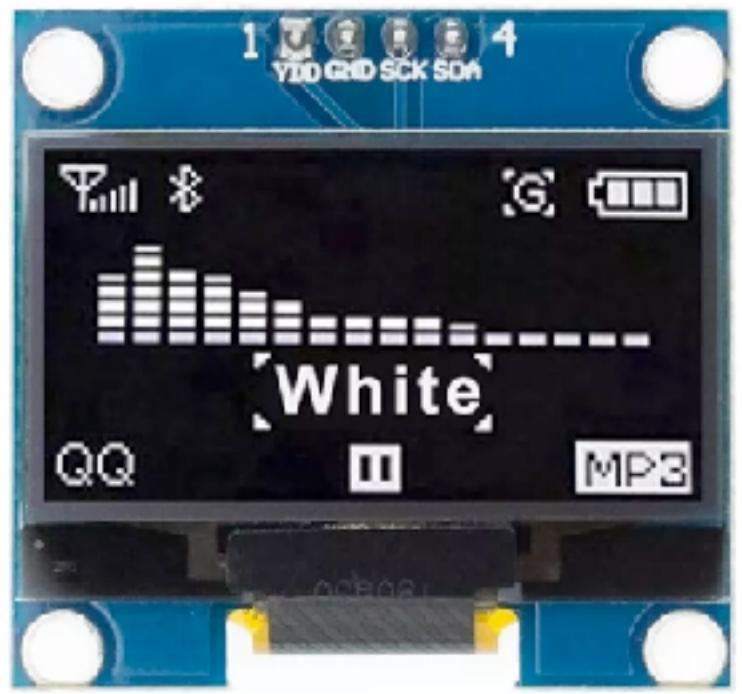
- **Example:** assume you want to drive an **OLED** display via **I²C**
- These are the **questions** you need to answer:
 - ▶ Does the chip natively support **I²C** or do we have to implement it *from scratch* (e.g, via bit-banging)?
 - ▶ Does the Arduino API already provide a general **driver** for **I²C**?
 - ▶ What is the *particular* **protocol** to communicate with the display?
 - ▶ Does a **library** for communication with the display already exists?
 - ▶ If so, does it fulfil our requirements in terms of **capabilities, resources, configurability**, and so on.



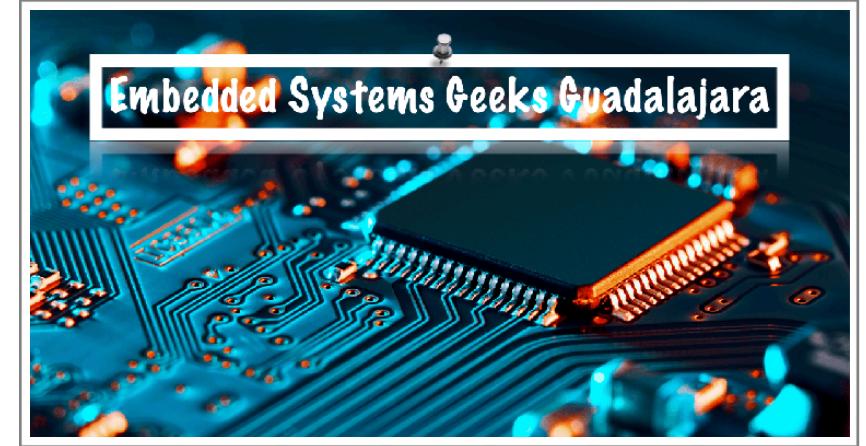
Knowing your Chip (cont'd)



- **Example:** assume you want to drive an **OLED** display via **I²C**
- These are the **questions** you need to answer:
 - ▶ Does the chip natively support **I²C** or do we have to implement it *from scratch* (e.g., via bit-banging)? → **YES**
 - ▶ Does the Arduino API already provide a general **driver** for **I²C**? → **YES**
 - ▶ What is the *particular* **protocol** to communicate with the display?
 - ▶ Uses the SSD1306 driver/controller chip ([SSD1306.pdf](#))
 - ▶ Does a **library** for communication with the display already exists? → **YES**
 - ▶ If so, does it fulfil our requirements in terms of **capabilities, resources, configurability**, and so on. → **YES/NO? DEPENDS ON PROJECT ...**

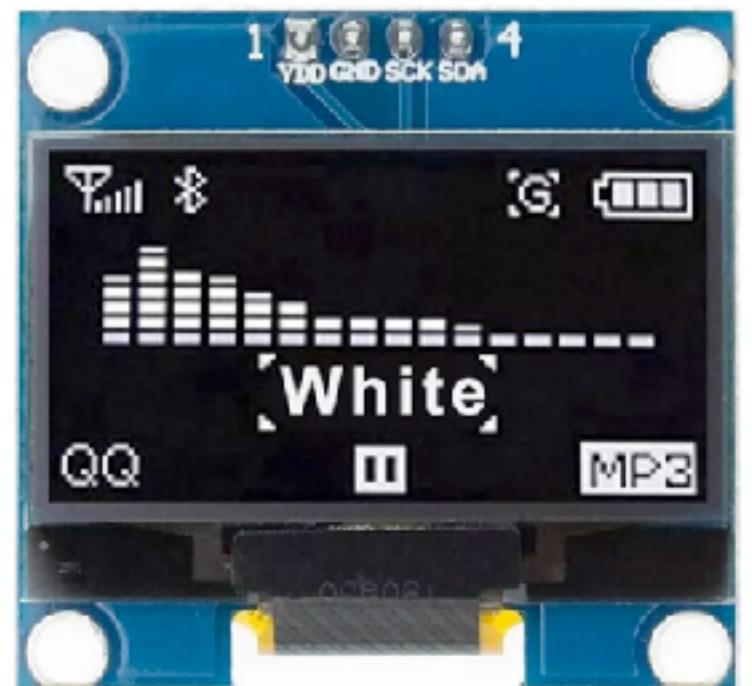


Knowing your Chip (cont'd)

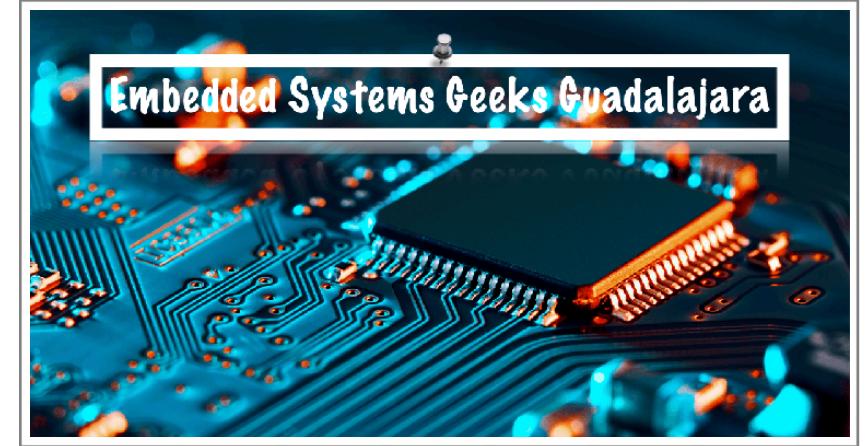


- **Advice:**

- ▶ Using a ready-made library can sometimes save us **time** and **trouble** to get something working (more) quickly.
- ▶ If we want a more **robust** / **specialised** / **maintainable** / **bespoke** driver, we **might** have to write it ourselves.
- ▶ Like with all third-party and open-source software, expect **bugs**.
- ▶ Embedded programming gives us the choice to either ...
 - ✓ work at a **higher** abstraction level via API / components
 - ✓ work at a **lower** abstraction level via direct I/O access
 - ✓ try to balance **effort** against **benefit** of reinventing stuff (!)

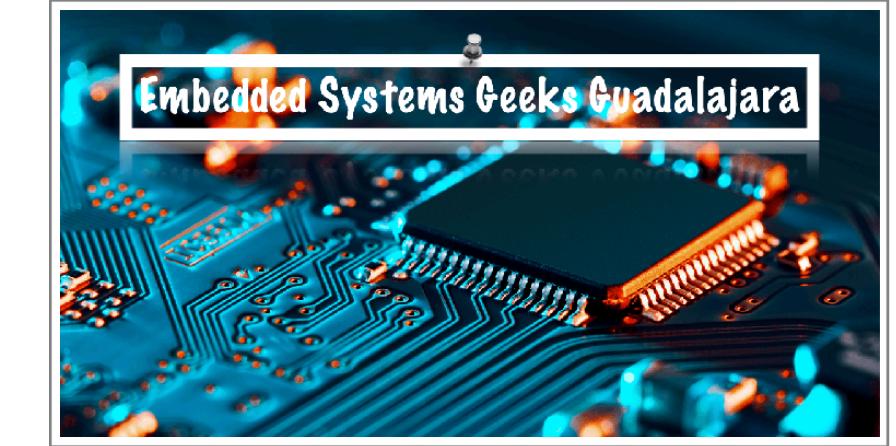


Example: OLED SSD1306 Driver



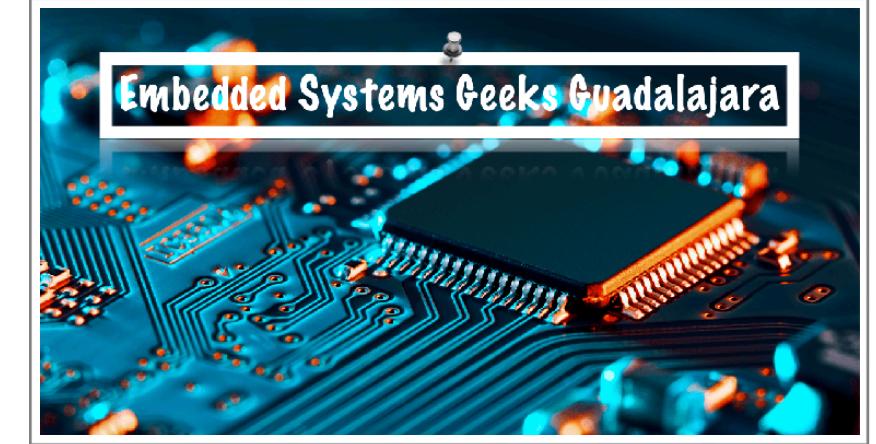
- Driver that I started to write with ([memory](#)) economy in mind.
- In most cases, we do **not** need a **frame buffer** ...
- Saves us **1 KB** of precious SRAM, i.e., for a 128x64 dots OLED.
- Also a good exercise to learn about **I²C** (more on this to follow)
- **Question:** When *do* we need a frame buffer?
 - ➔ Device does **not hold** display data by itself (needs refresh)
 - ➔ Data can only be written in **words** (not individual **bits**) ...
... and we require to draw objects on top of each other!
 - ➔ We need to **read out** display data for one reason or another
(e.g., to save the image in a paint program, hibernate device, etc.)

Communication Protocols

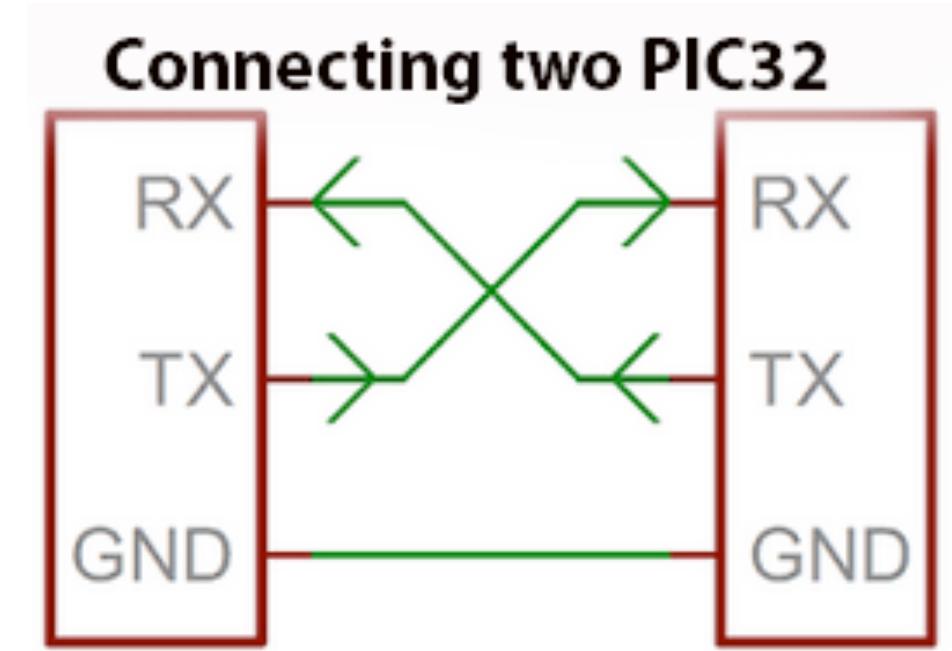


- How do protocols *in general* differ from each other?
 - **analogue** (modulated) vs **digital** (low / high)
 - **synchronous** (**clocked**) vs **asynchronous** (**self-clocking**)
 - permissible data **transmission rate** (**frequency**)
 - number of **signals** / **wires** required; support for **duplex**?
 - number of **senders** (**masters**) and **receivers** (**slaves**) allowed
 - built-in **error checking** mechanisms?
 - **electrical characteristics** (**voltage**, **sensitivity to noise**, and so on)
 - underlying **protocol specification** and **standardisation**

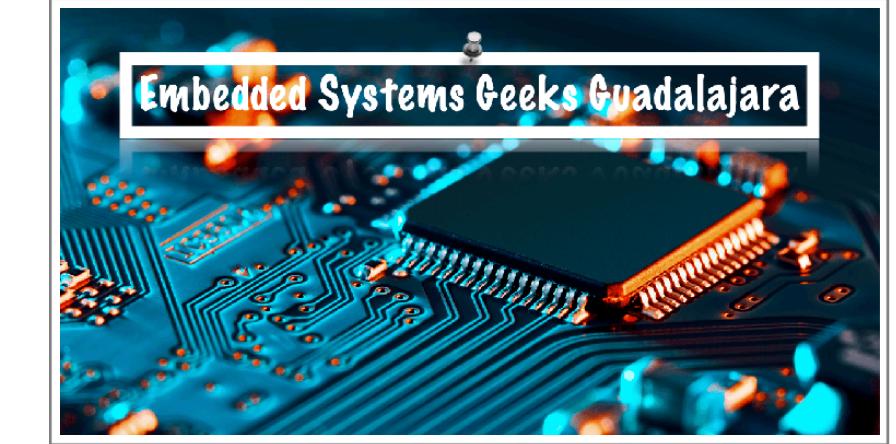
Protocol Example: UART



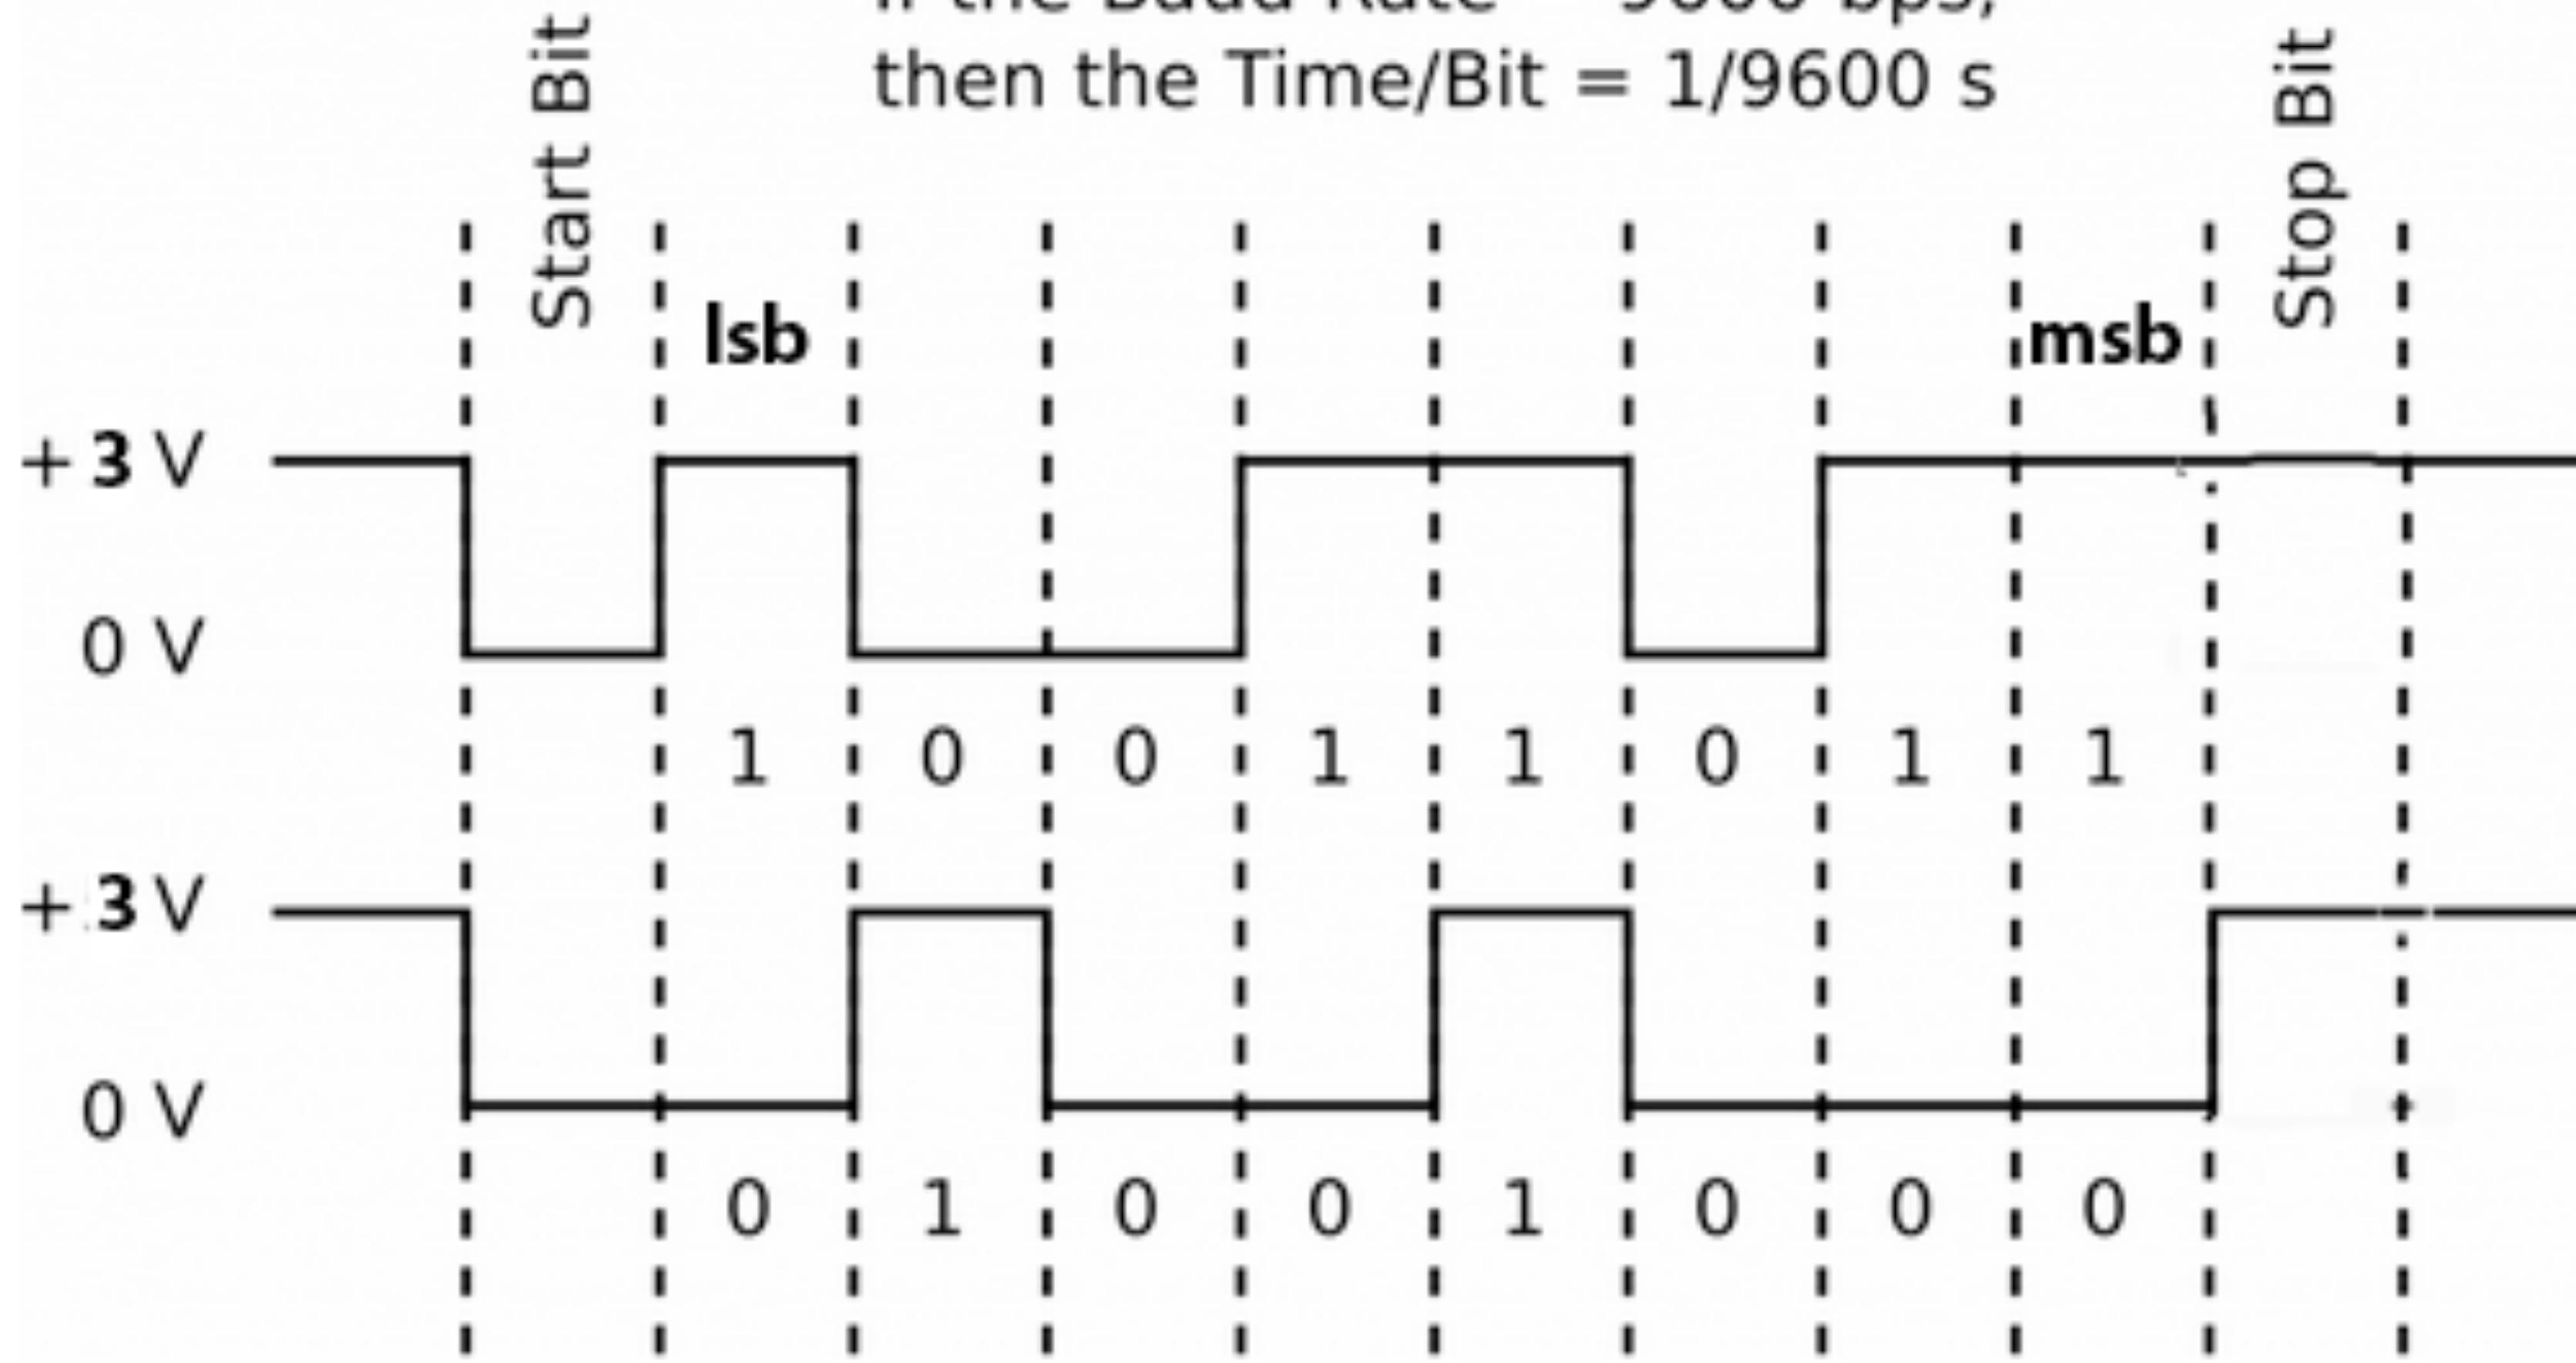
- **UART** = Universal **A**synchronous **R**eceiver-**T**ransmitter
- Requires two wires for full bidirectional transmission:
 - ➔ **TX** (transmit) and **RX** (receive), plus **GND** (ground wire).
- **Asynchronous**, using a fixed predetermined BAUD rate.
- Typical speeds are **9,600, 19200, 38400, 57600, and 115200 bps**.
- Connects **two devices**, no multi-master or multi-slave by design.
- Supports error checking via a **parity bit**.
- Supported natively by the Arduino / AVR family (!)
 - ➔ Actually used, AFAIK, to upload programs via the Arduino boot-loader ...



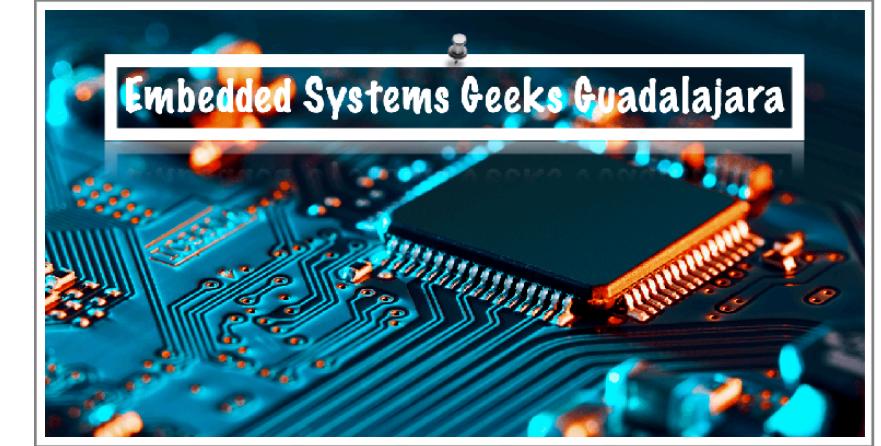
Protocol Example: UART



If the Baud Rate = 9600 bps,
then the Time/Bit = $1/9600$ s

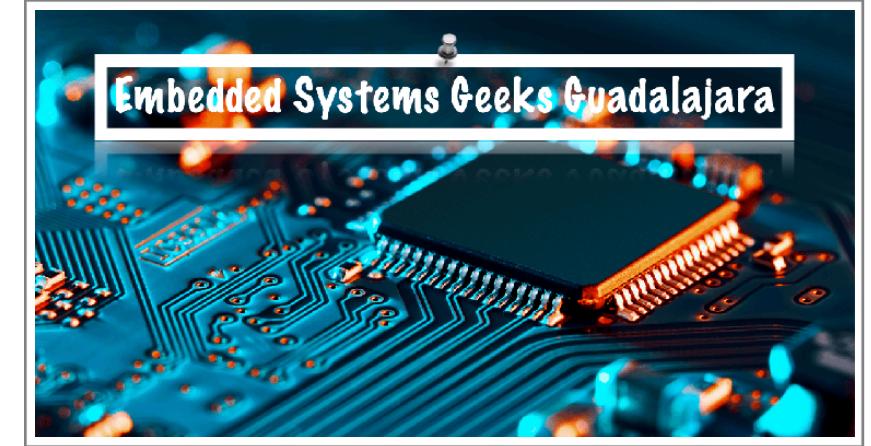


Protocol Example: I²C

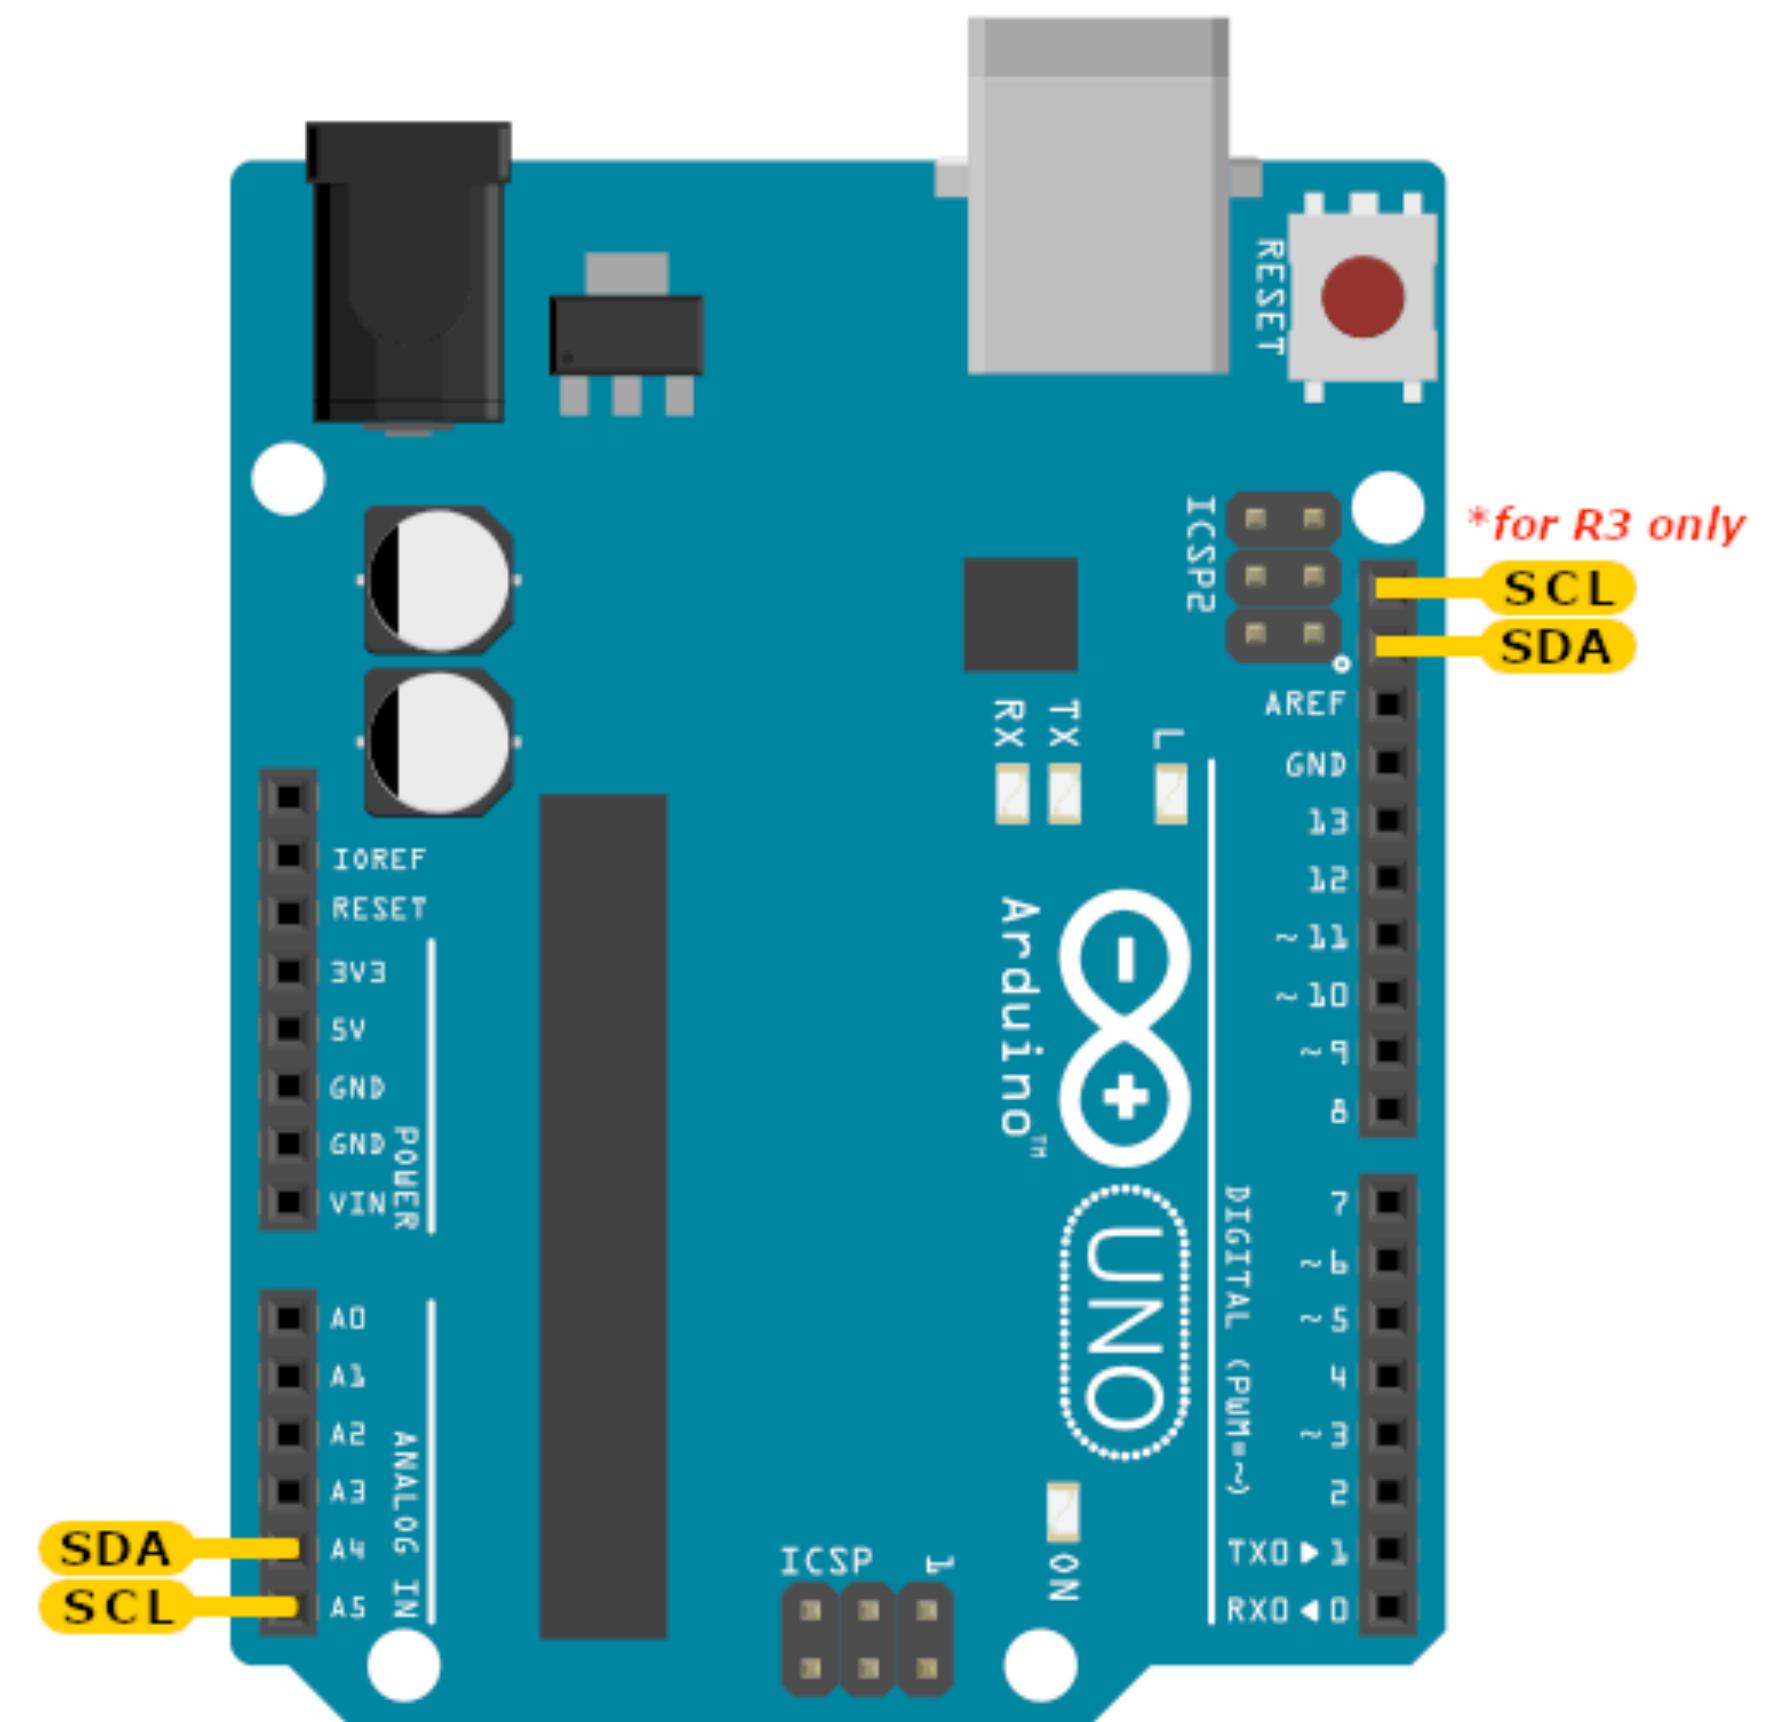
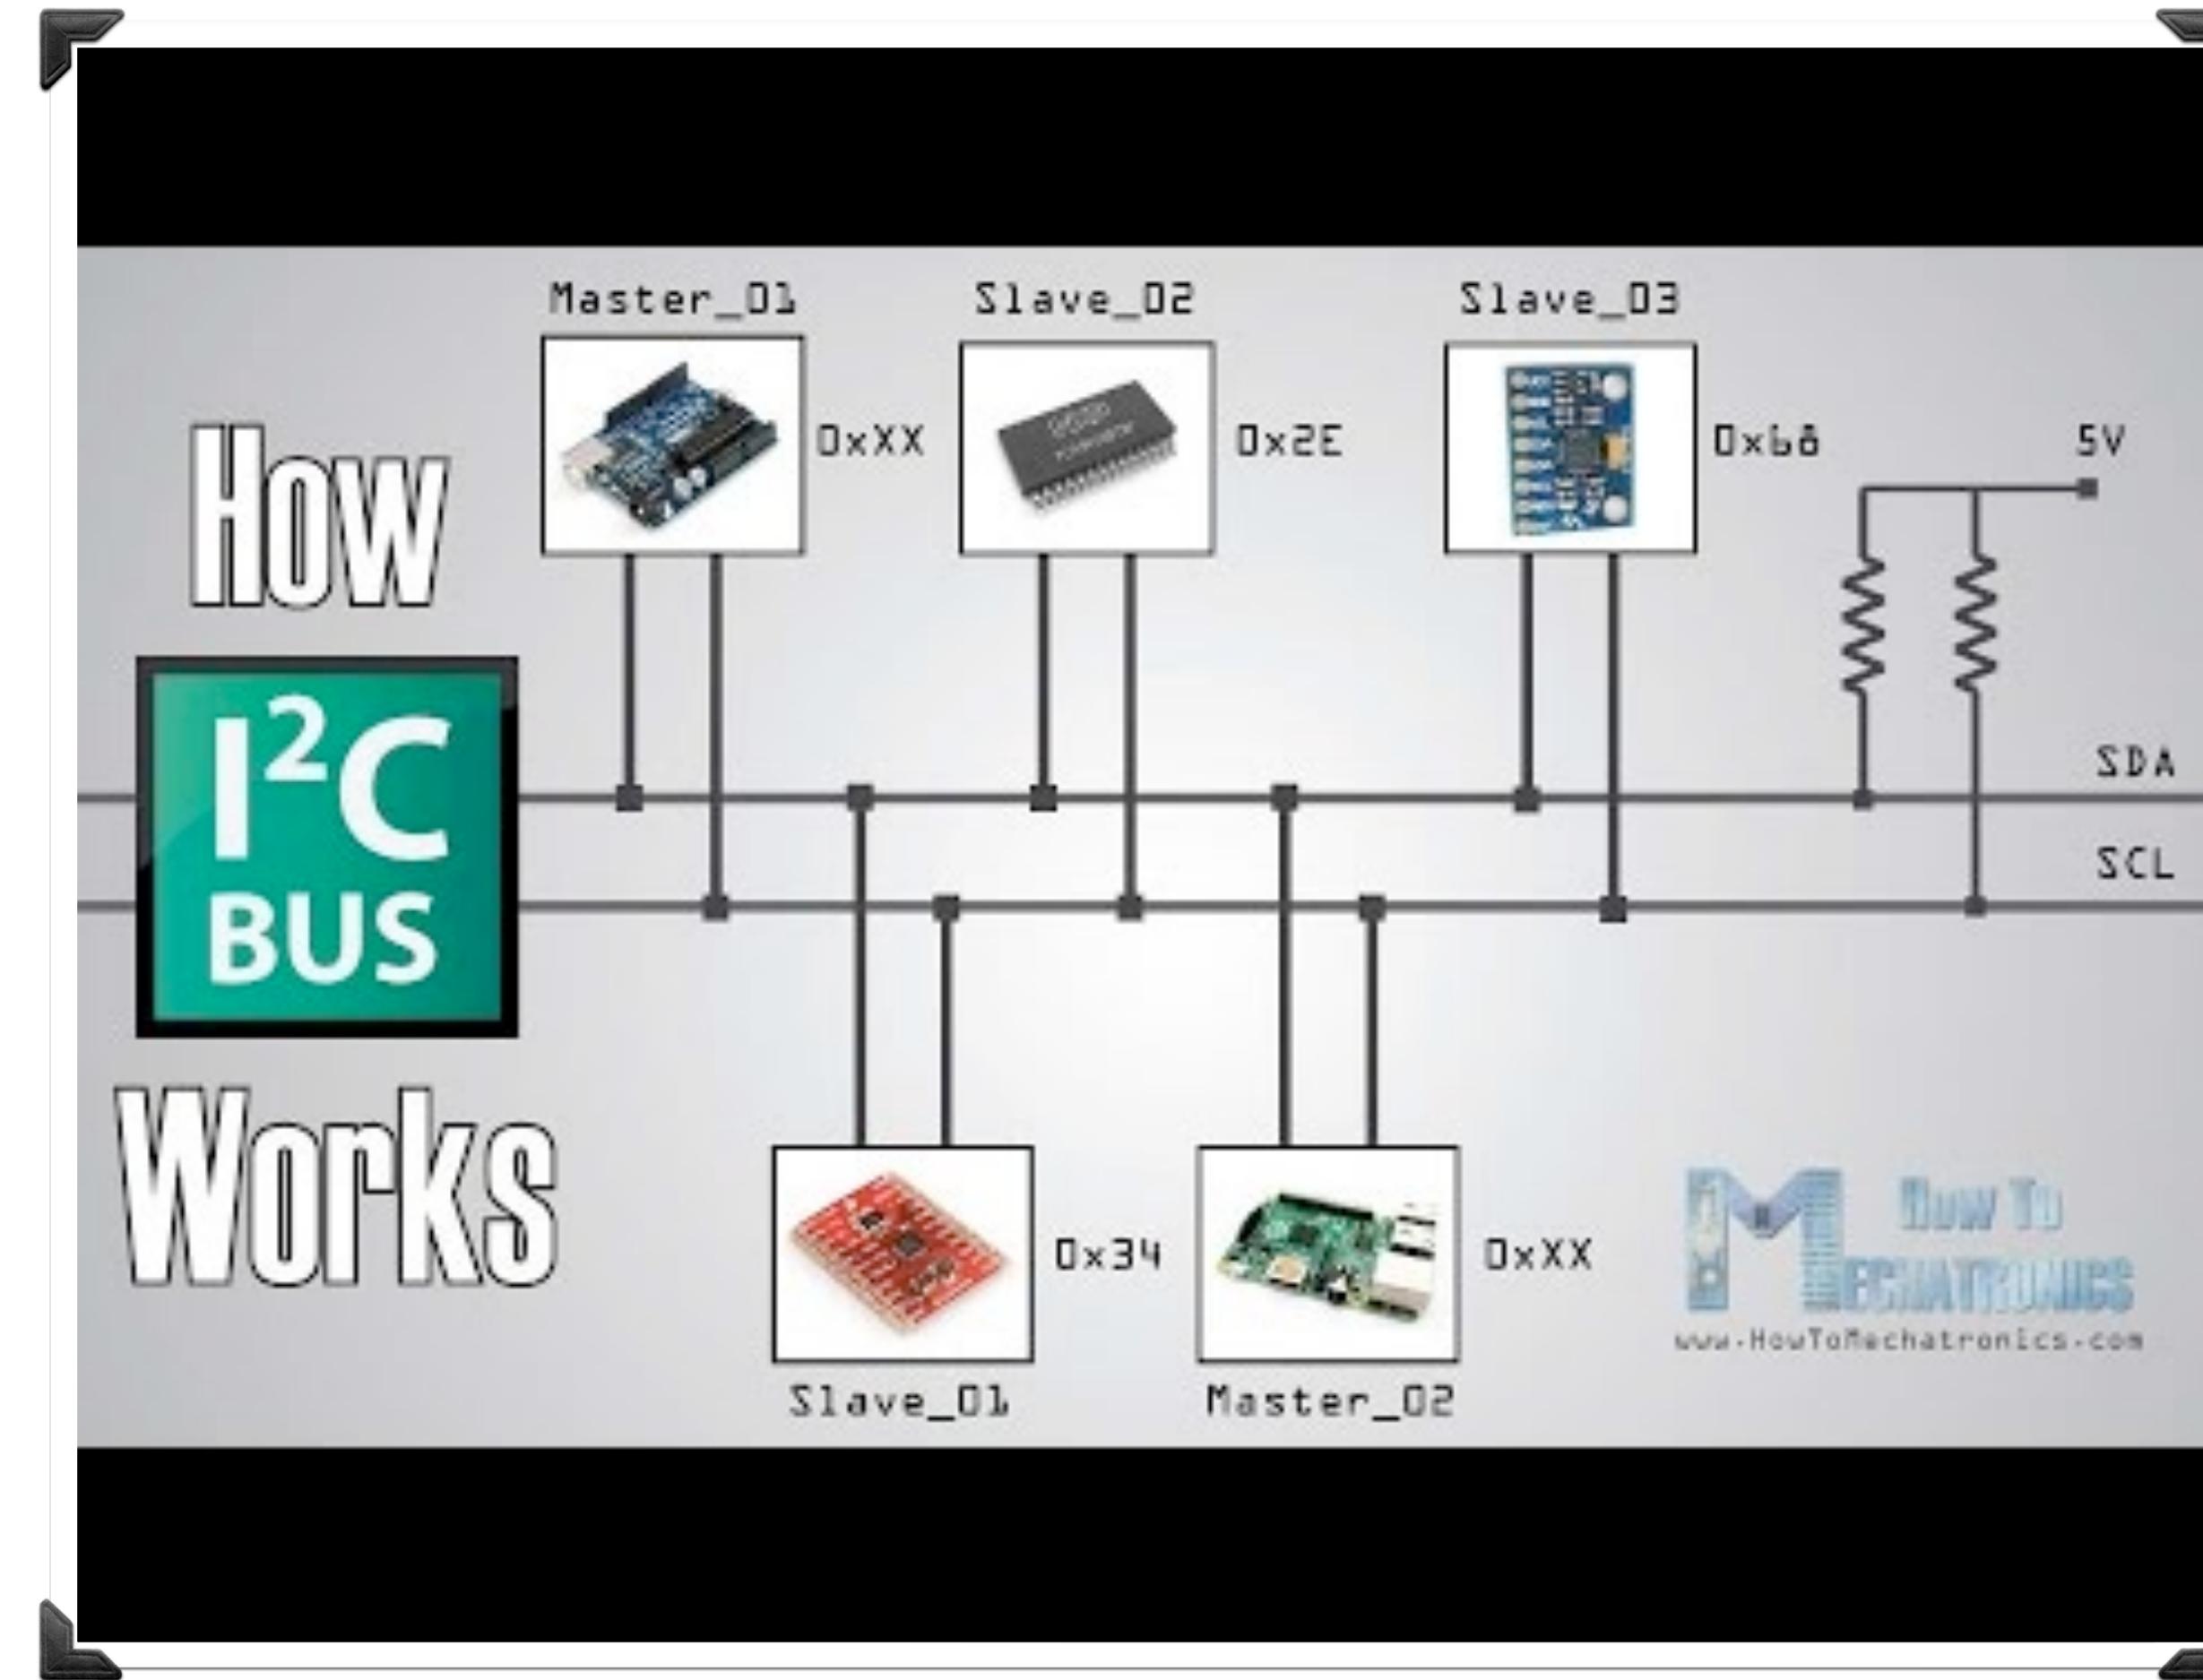


- Supports **multiple masters** and **multiple slaves**.
- Requires two wires (**SDA** and **SCL**) plus GND (ground).
- Each slave can be talked to individually via its *address*.
- Masters can send at the same time (outputs use open-drain).
 - To output a 1 go into **tristate** (high impedance)
 - To output a 0 set to **GND**
 - Requires a **pull-up resistor** (the MCU already provides this)
- Supported speeds are: **100 kbit/s (standard mode)**, up to **400 kbit/s (fast mode)**, and up to **1 Mbit/s (fast mode plus)**.
- Many components speak I²C (memories, ADC/DAC, RGB LEDS, ...)

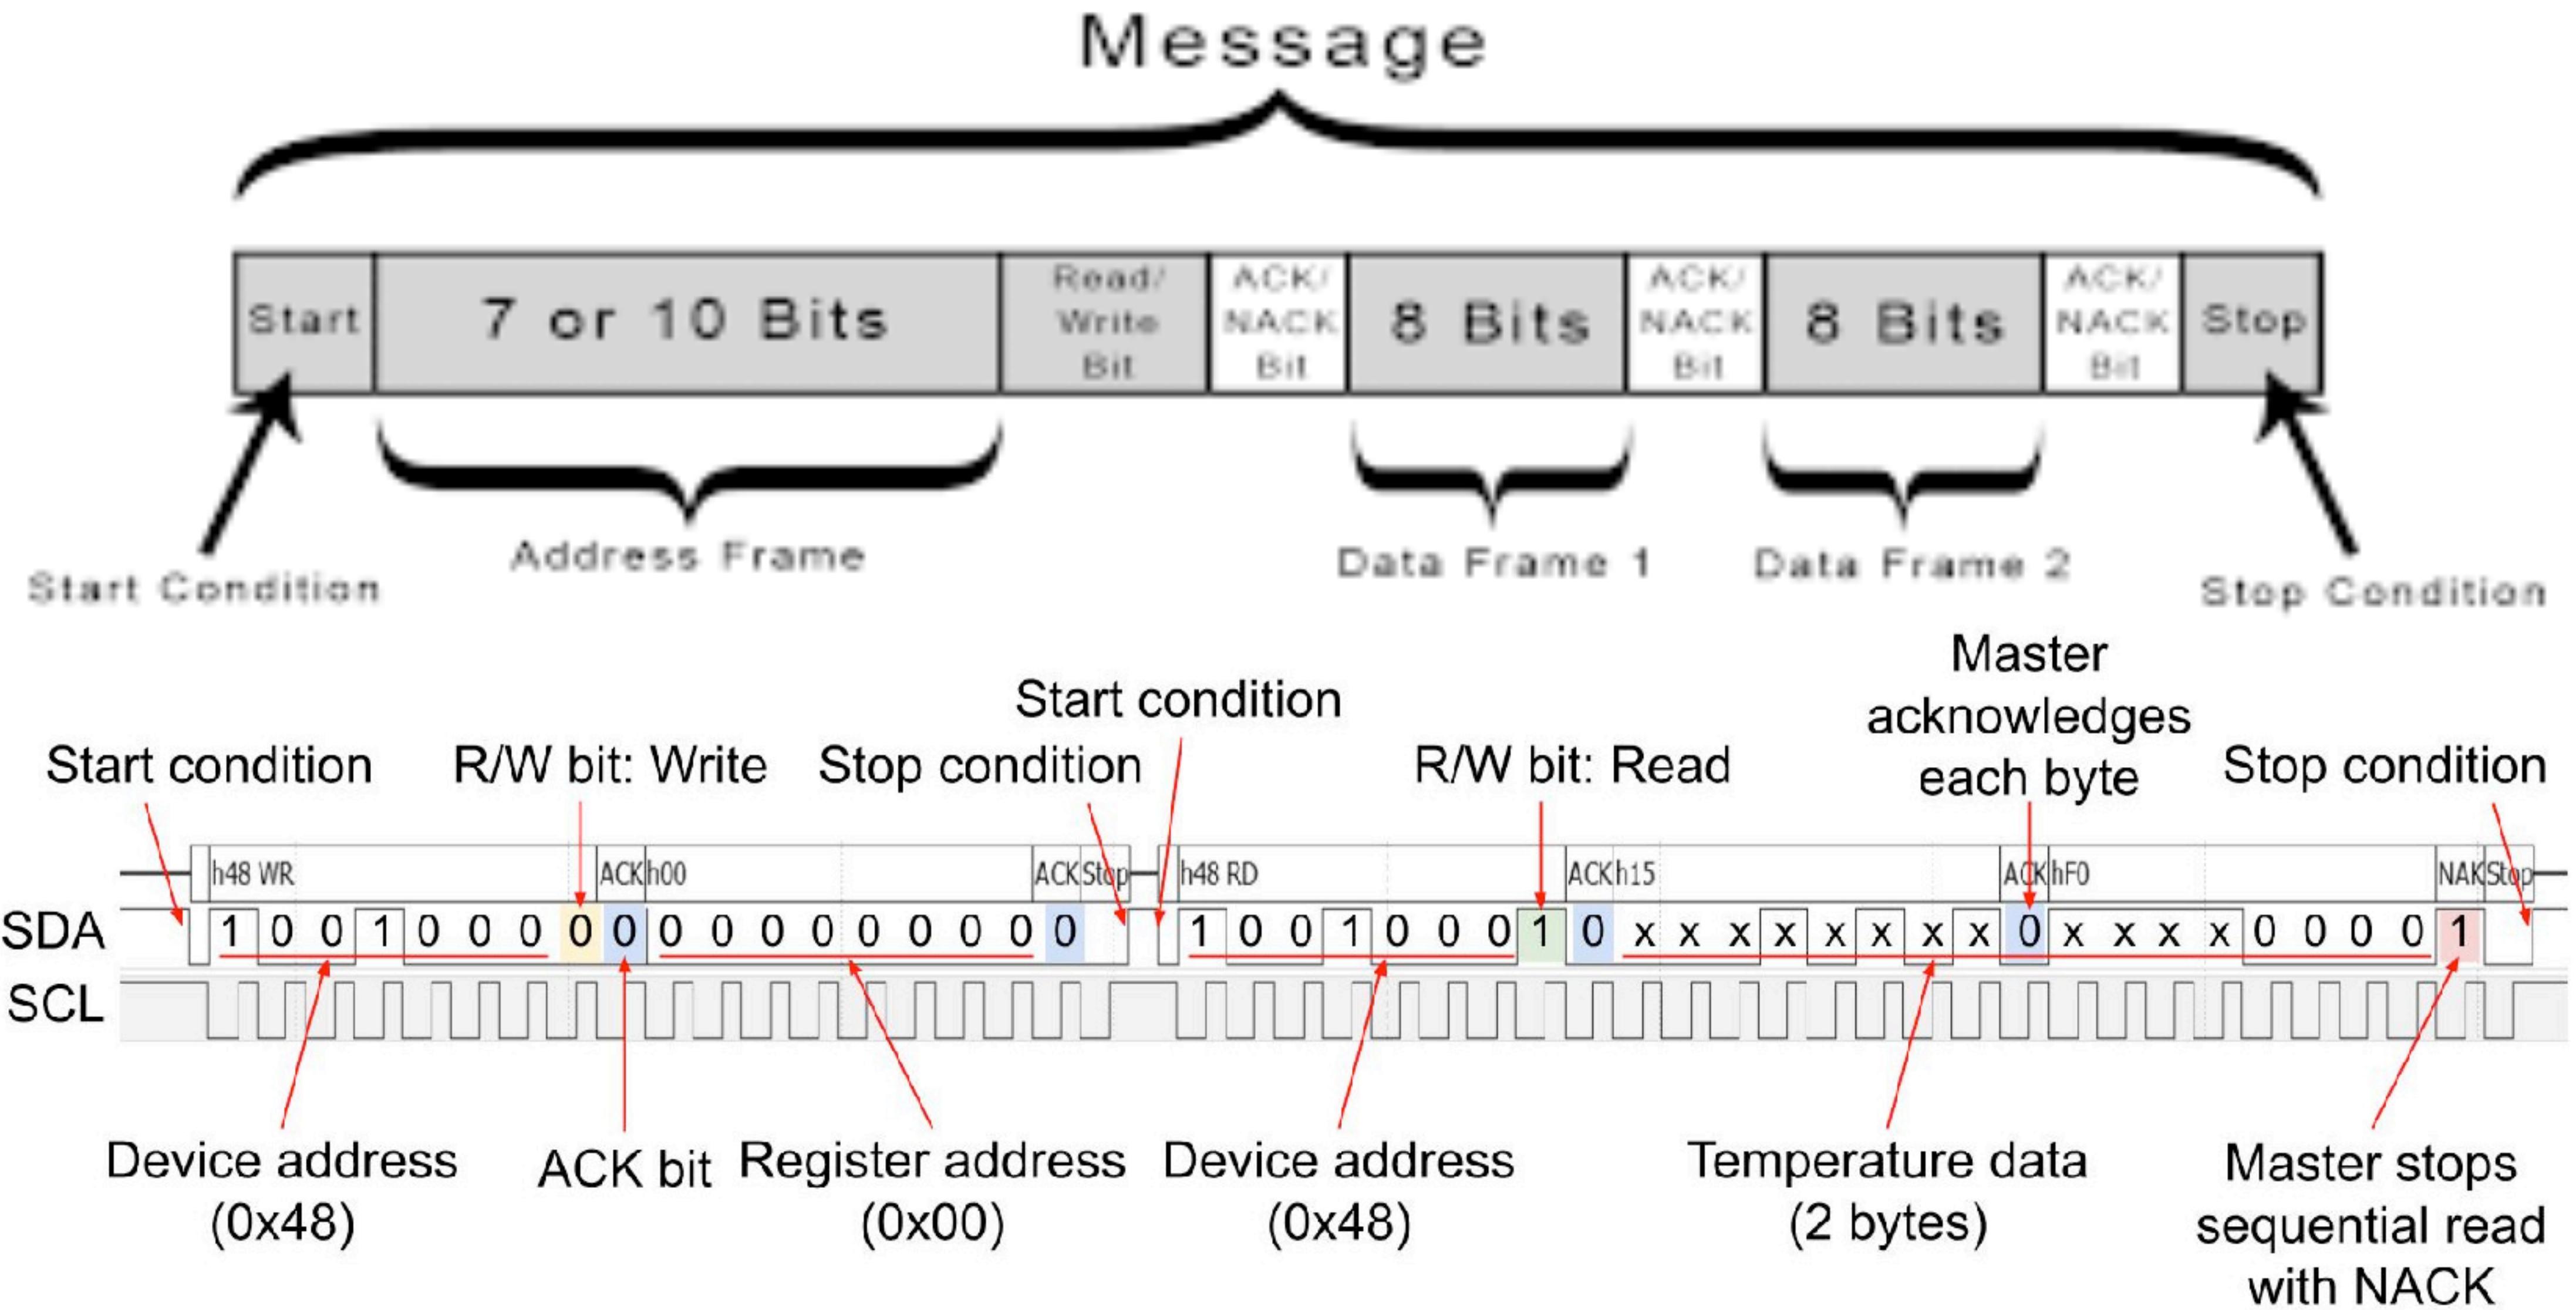
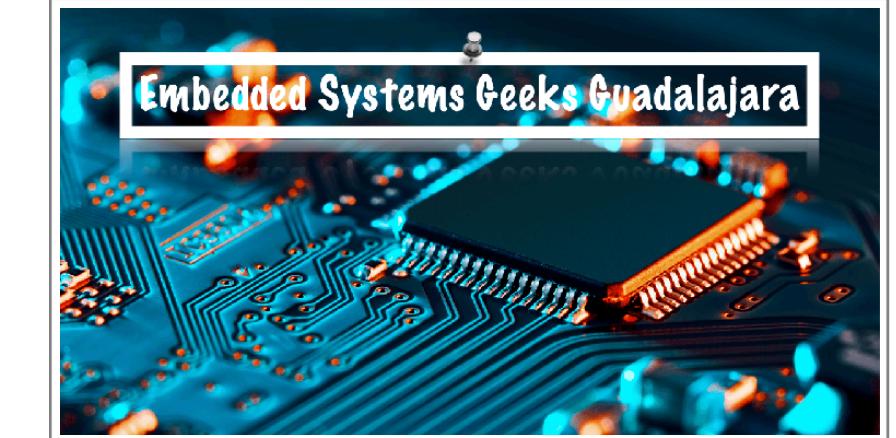
Protocol Example: I²C



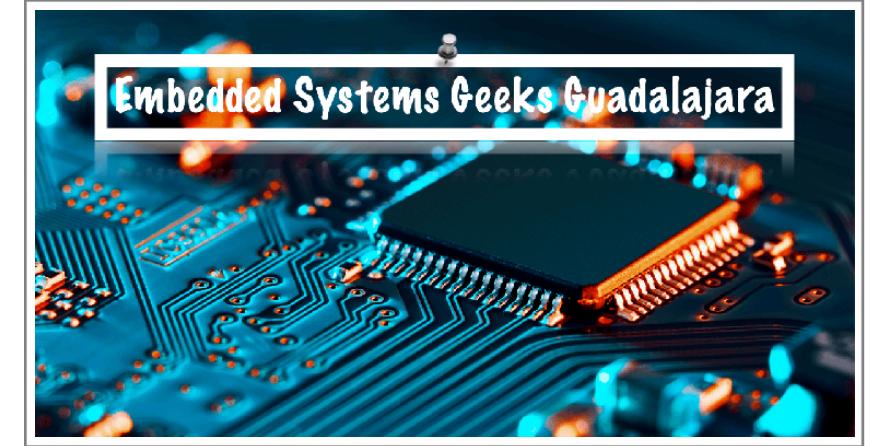
- Many devices can be connected to SDA and SCL.



Protocol Example: I²C

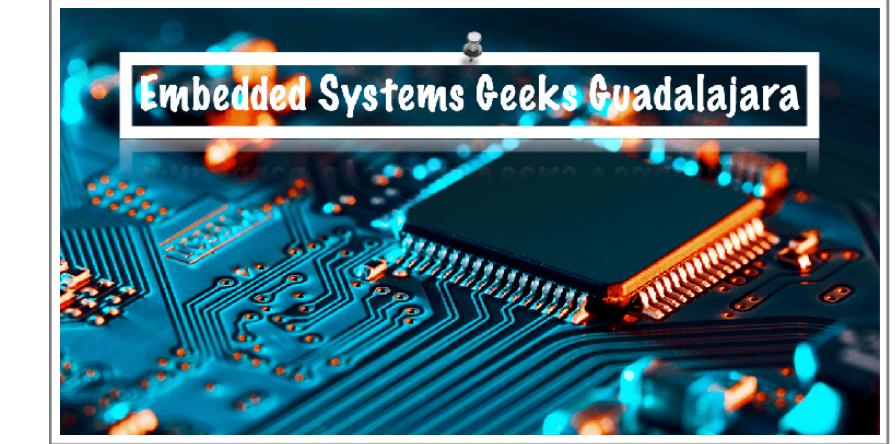


I²C on Arduino



- Supported by the **I²C / TWI** library (Wire.h).
- Takes advantage of the native chip support, using interrupts.
 - ➔ But there might be some limitations that require **busy waiting** ...
- Programmer just has to focus on the transmitted & received data words, and not any details related to the hardware protocol.
- Implementing **I²C** in a fully complaint manners it no easy, e.g.,
 - ➔ How to *detect and deal with* **multiple masters** sending **at once**?
- Using the correct PIN output setting is crucial (**open drain** with **pull-up**).
- **I²C** is usually only used for on-board hardware (short connections)

I²C on Arduino (cont'd)



- Arduino **Wire** library code example:

```
13 #include <Wire.h>
14
15 void setup()
16 {
17     Wire.begin(); // join i2c bus (address optional for
18 }
19
20 byte x = 0;
21
22 void loop()
23 {
24     Wire.beginTransmission(4); // transmit to device #4
25     Wire.write("x is ");           // sends five bytes
26     Wire.write(x);              // sends one byte
27     Wire.endTransmission();    // stop transmitting
28
29     x++;
30     delay(500);
31 }
```

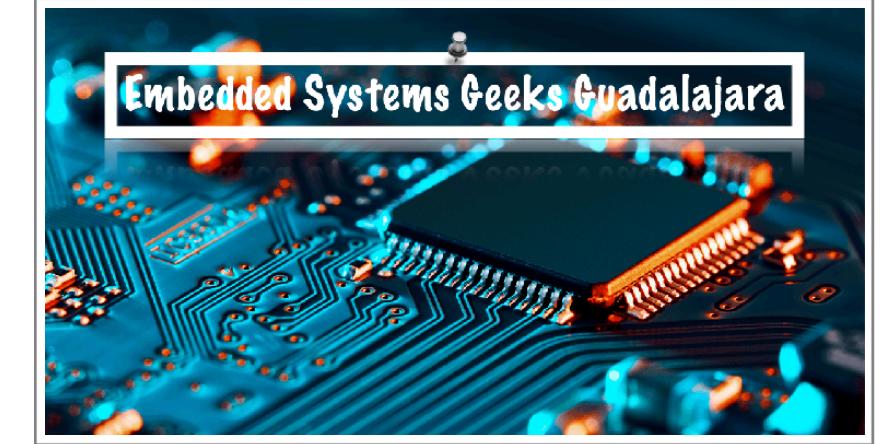
<https://docs.arduino.cc/learn/communication/wire>

Example Revisited: SSD1306 Driver



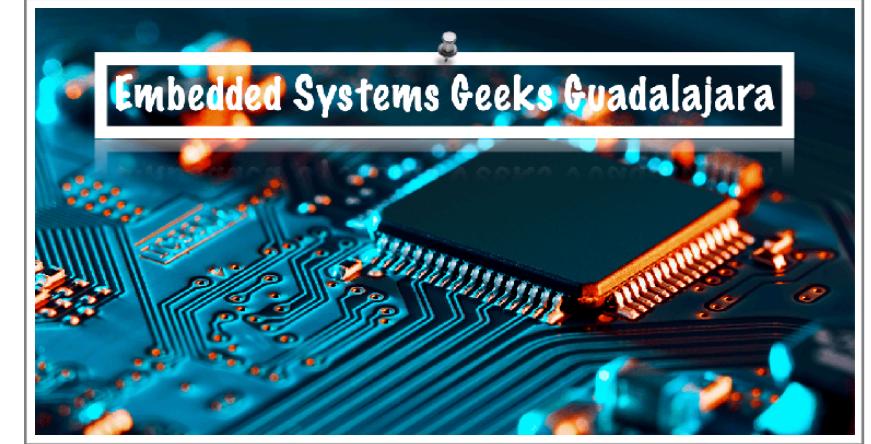
- Have a quick look at the source code in the **Arduino IDE** if time ...

Protocols: Summary



- We have only discussed **two** common protocols here, but there exist **many more** ...
- **Recommendation:** look into them when you have a **concrete** need !
- Most of them are **not** difficult to understand, but others may be *a little more challenging*. E.g.,
 - Composite analogue video signals (**PAL, NTSC, SECAM**)
 - Other graphics signal standards (**VGA, DVI, HDMI**)
- You might reach the limits of the Arduino MCUs with that ...
 - **Linus Åkesson** did some cool stuff related to this:
<http://www.linusakesson.net/scene/phasor/index.php>

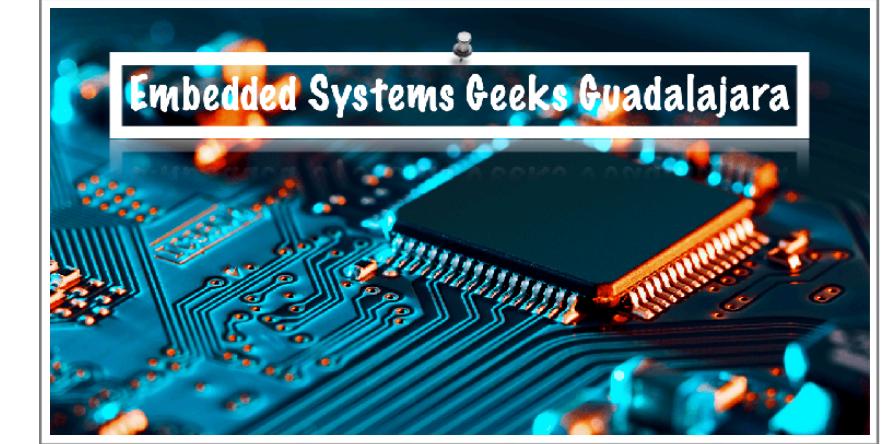
Misc Topic: Timing



- **Timing** is rather important when programming for embedded systems:
 - Devices often need to **synchronise** with each other.
 - Communication **protocols** have hard timing constraints.
 - Delays are sometimes implemented in pure software.
 - Watchdogs need to be triggered to keep the system alive.
 - And *unresponsive system* often means **death** ...
 - Data that is **sent** or **arrives late** may not be of use anymore ...
- Games development in the 80s:
 - Count CPU clock cycles to “**race the beam**”.

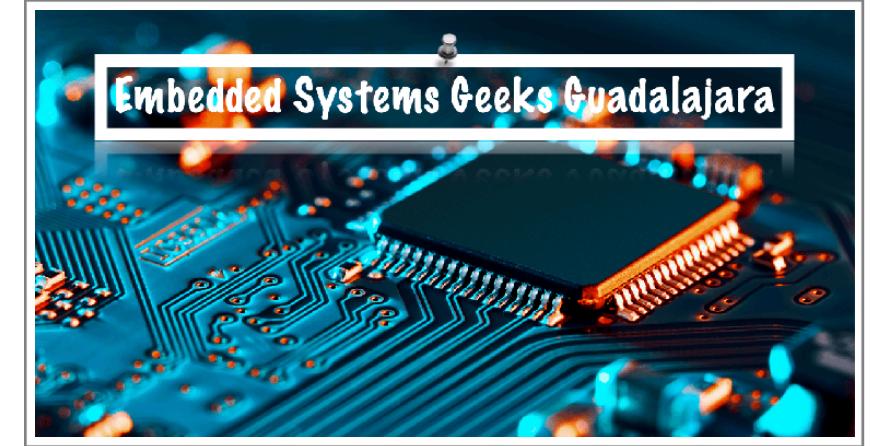


Misc Topic: Power/Load



- Electrical devices consumer power when they are in operation.
- Power is determined by **both** voltage and current: $P = U*I$
- **Problem:** output PINs of an MCU can only support current flow **up to a certain threshold**.
 - ▶ Anything above this can **damage / destroy** the chip 
- Equally, connecting an output to GND and writing a **high value** (1) can cause instant death ... (depends on internal safe-guarding mechanisms)
- Hence, always check that the **consumer** do not draw more current than permitted. Outputs must **never** be connected **directly** to V_{cc} or **GND**.
- Where do we find this information? In the **data sheet** !

DC/AC Characteristics

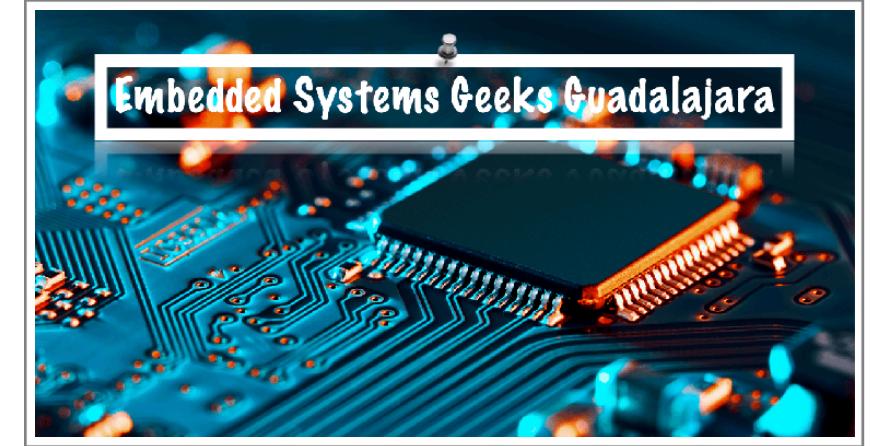


- Maximum Absolute Ratings for **ATmega328P** (page 258)

| Parameters | Min. | Typ. | Max. | Unit |
|---|------|-----------------|----------------|------|
| Operating temperature | -55 | | +125 | °C |
| Storage temperature | -65 | | +150 | °C |
| Voltage on any pin except <u>RESET</u> with respect to ground | -0.5 | | $V_{CC} + 0.5$ | V |
| Voltage on <u>RESET</u> with respect to ground | -0.5 | | +13.0 | V |
| Maximum operating voltage | | 6.0 | | V |
| DC current per I/O pin | | 40.0 | | mA |
| DC current V_{CC} and GND pins | | 200.0 | | mA |
| Injection current at $V_{CC} = 0V$ | | $\pm 5.0^{(1)}$ | | mA |
| Injection current at $V_{CC} = 5V$ | | ± 1.0 | | mA |

Note: 1. Maximum current per port = $\pm 30mA$

DC/AC Characteristics



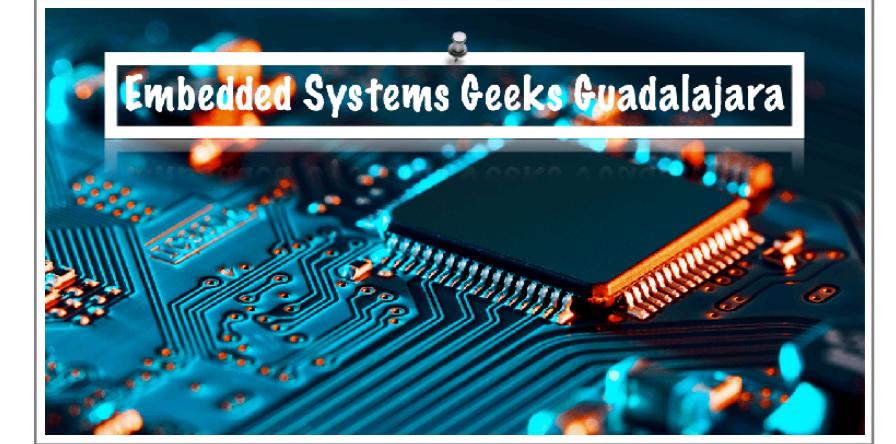
- Maximum Absolute Ratings for **ATmega328P** (page 258)

| Parameters | Min. | Typ. | Max. | Unit |
|---|------|-----------------|----------------|------|
| Operating temperature | -55 | | +125 | °C |
| Storage temperature | -65 | | +150 | °C |
| Voltage on any pin except <u>RESET</u> with respect to ground | -0.5 | | $V_{CC} + 0.5$ | V |
| Voltage on <u>RESET</u> with respect to ground | -0.5 | | +13.0 | V |
| Maximum operating voltage | | 6.0 | | V |
| DC current per I/O pin | | 40.0 | | mA |
| DC current V_{CC} and GND pins | | 200.0 | | mA |
| Injection current at $V_{CC} = 0V$ | | $\pm 5.0^{(1)}$ | | mA |
| Injection current at $V_{CC} = 5V$ | | ± 1.0 | | mA |

Note: 1. Maximum current per port = $\pm 30mA$

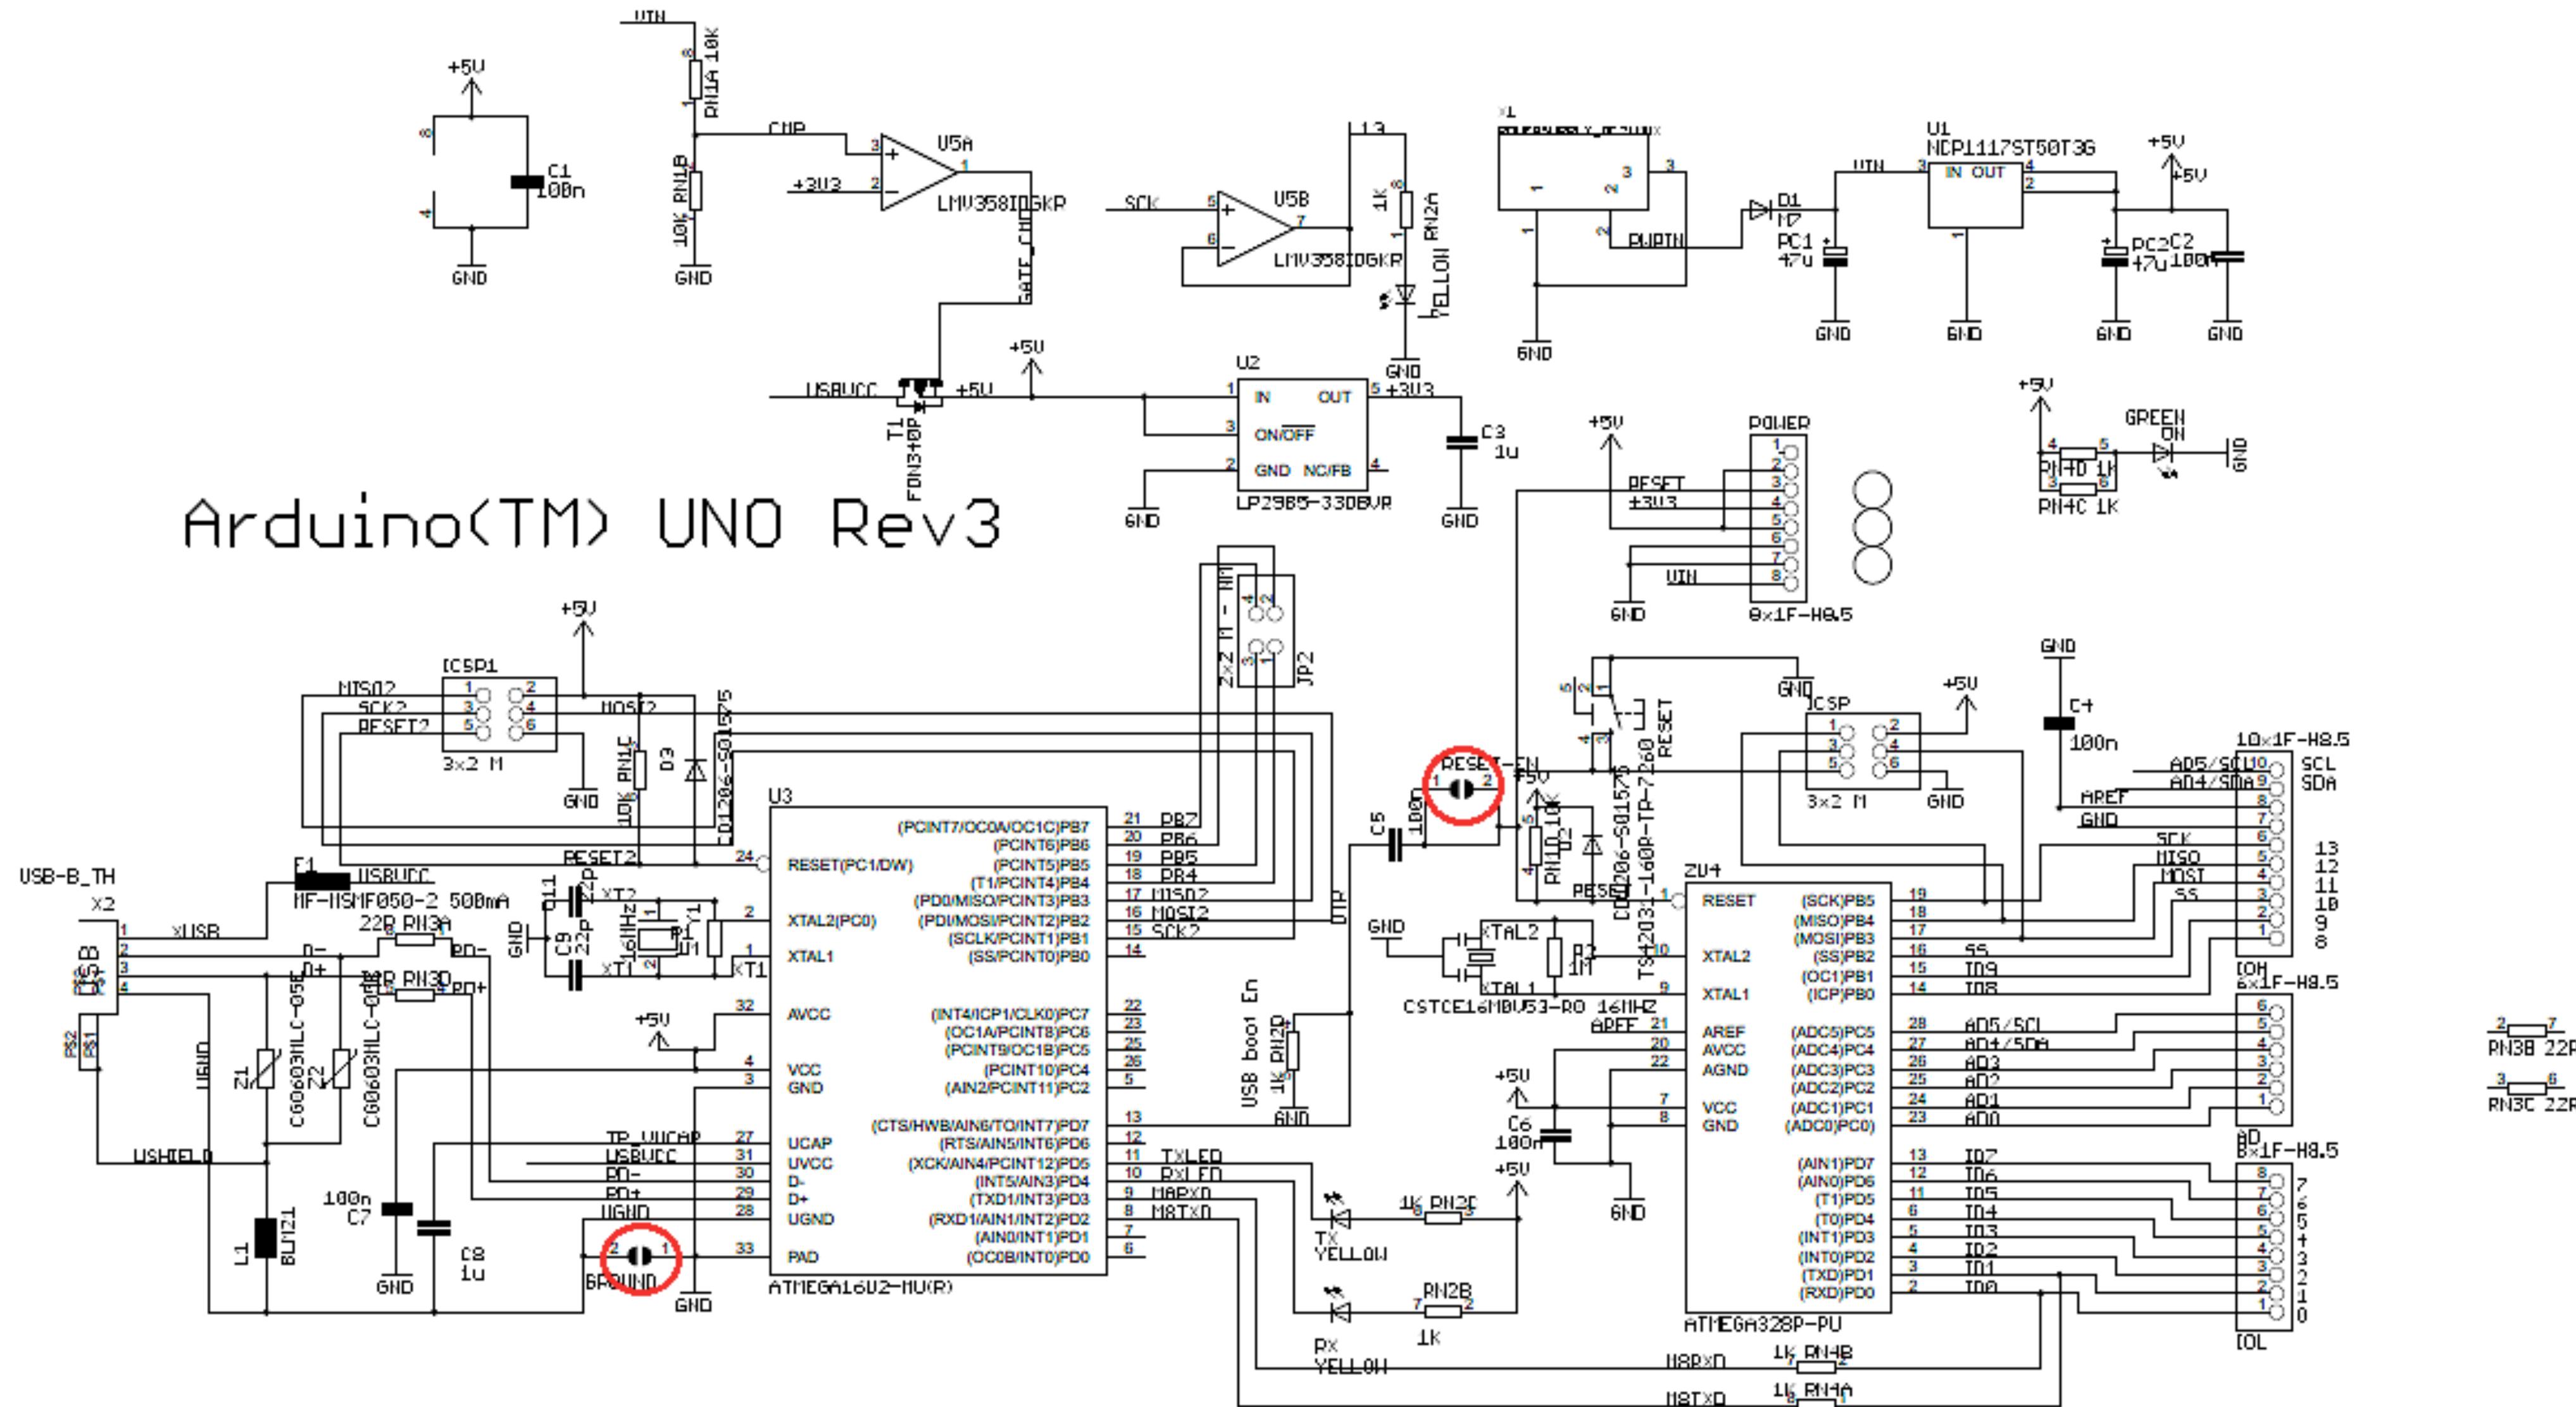
Max. permissible load on I/O PINs

Misc Topic: Schematics

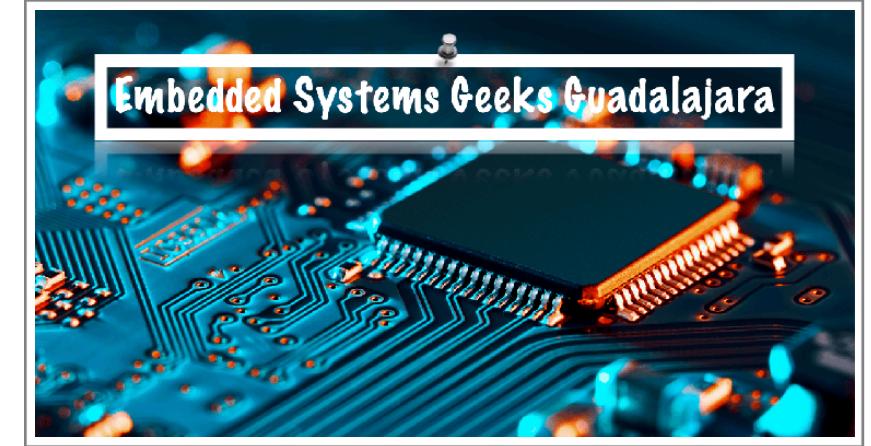


- To **really understand** a development **board or device**, we do not get around reading semantics.
- While an MCU includes *most things we need* for deploying *simple* embedded applications, we often profit from **additional components**.
- Example: USB → UART bridge (**FTDI**), **graphics** signal generators, Ethernet **PHY**, WiFi **transceivers**, MOSFET **drivers** for motors, etc.
- Again, not **rocket science** 🚀 to understand them, but requires basic knowledge in electrical **symbols** and **conventions** of circuit design.

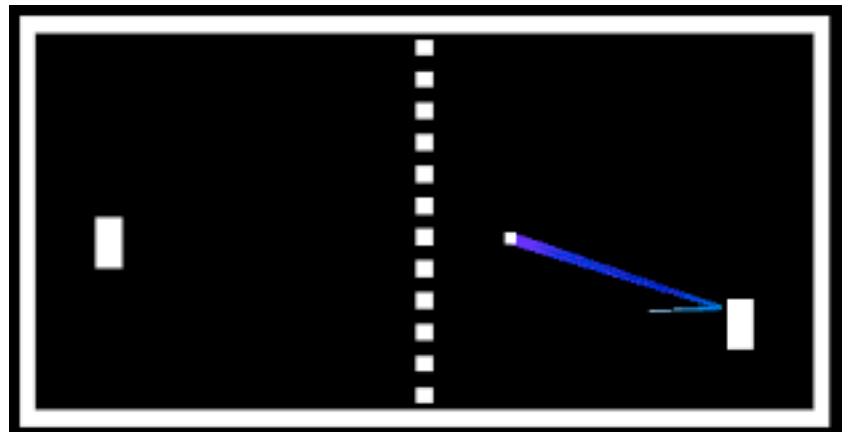
Arduino UNO Schematics



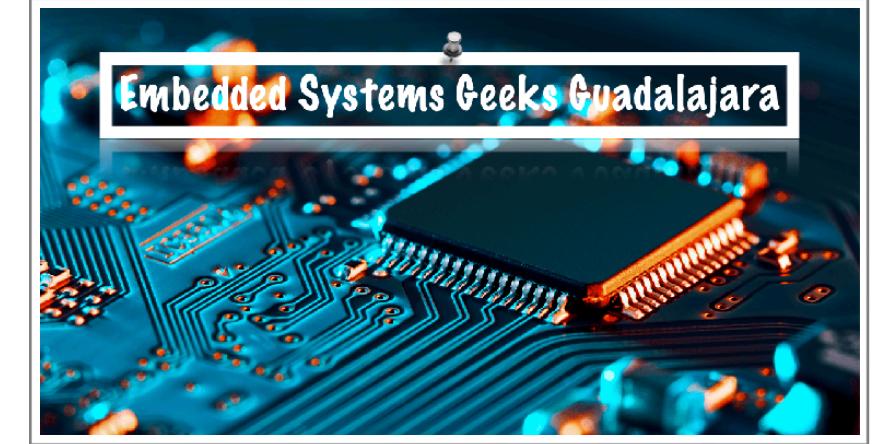
Conclusion



- We had a brief look at the Arduino and AVR **architecture**.
- We understood *a little more* the **constraints** when developing for MCU.
- We examine several hardware communication **protocols**.
- We discussed several miscellaneous topics (**timing, load, schematics**)
- This should give you a **good start** to embark **on your own quest**:
 - ➔ We cannot discuss everything but there is a lot of material on the **WEB**.
 - ➔ Hardware engineering and programming is not rocket science ...
 - ➔ ... but it requires a certain shift in **thinking** and **paradigm**.
 - ➔ A suggestion: Let's develop an **arcade game** together !



Further Reading & Study



- **First stop:** **Documentation** on <https://www.arduino.cc/>
 - Language comparison: Arduino *Wiring* Language vs plain C/C++
<https://docs.arduino.cc/tutorials/generic/language-comparison>
 - Arduino Language Reference which can be found on the above website.
- **Next:** Try out some of the examples that come **preinstalled** with the Arduino IDE: <https://docs.arduino.cc/library-examples/>
- Have a brief look at the **data sheet** of the [ATmega38P](#)
- [Arduino Wiki](#) of **NUS** (National University of Singapore)
- [Arduino Cookbook](#) (especially, if you are new to programming)
- Any **good book** on C/C++ programming

Thanks for Listening!



Time for **questions** and **comments**.

